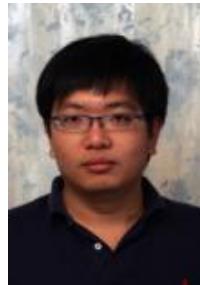




A New Parallel Framework for Machine Learning

Joseph Gonzalez

Joint work with



Yucheng
Low



Aapo
Kyrola



Danny
Bickson



Carlos
Guestrin



Alex
Smola



Guy
Blelloch



Joe
Hellerstein



David
O'Hallaron

Select Lab

Carnegie Mellon

How will we
design and implement
parallel learning systems?

We could use

Threads, Locks, & Messages

“low level parallel primitives”

Threads, Locks, and Messages

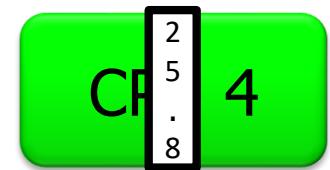
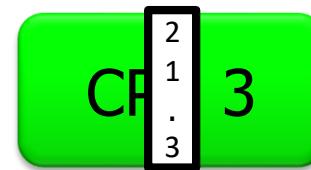
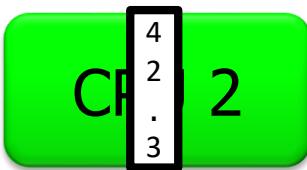
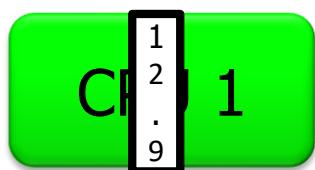
- Graduate students **repeatedly** solve the same parallel design challenges:
 - Implement and debug complex parallel system
 - Tune for a specific parallel platform
 - Two months later the conference paper contains:
“We implemented _____ in parallel.”
- The resulting code:
 - is difficult to maintain
 - is difficult to extend
 - couples learning model to parallel implementation

... a better answer:

Map-Reduce / Hadoop

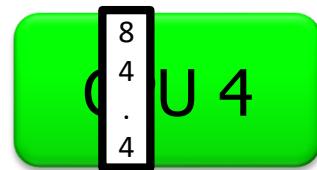
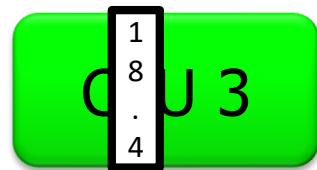
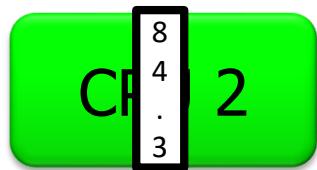
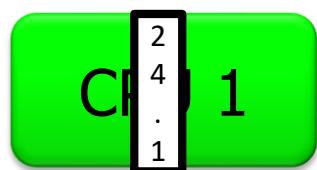
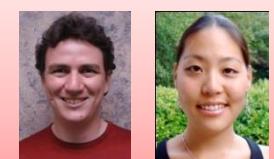
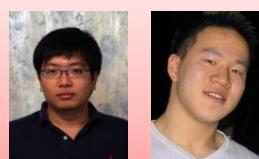
Build learning algorithms on-top of
high-level parallel abstractions

MapReduce – Map Phase



Embarrassingly Parallel independent computation
No Communication needed

MapReduce – Map Phase



1
2
.
9

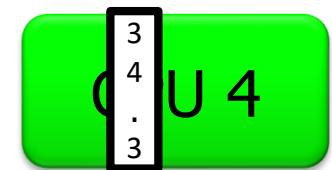
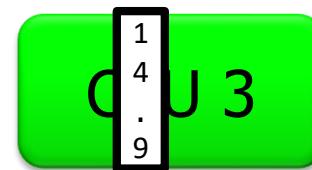
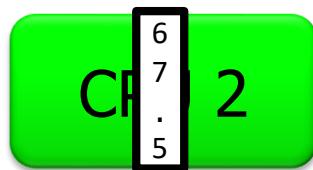
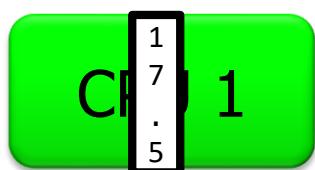
4
2
.
3

2
1
.
3

2
5
.
8

Image Features

MapReduce – Map Phase



1
2
.
9

2
4
.
1

4
2
.
3

8
4
.
3

2
1
.
3

1
8
.
4

2
5
.
8

8
4
.
4

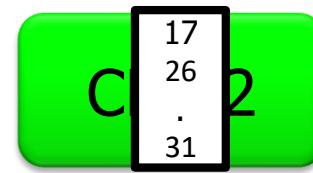
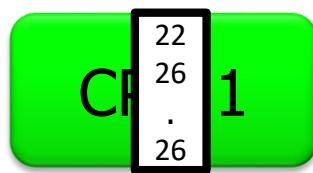
Embarrassingly Parallel independent computation

No Communication needed

MapReduce – Reduce Phase

Attractive Face
Statistics

Ugly Face
Statistics

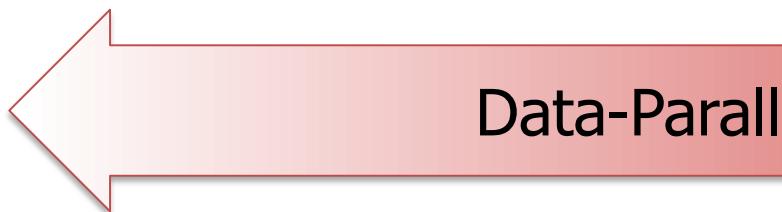


1 2 .9	2 4 .1	1 7 .5	4 2 .3	8 4 .3	6 7 .5	2 1 .3	1 8 .4	1 4 .9	2 5 .8	8 4 .4	3 4 .3
--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------

Image Features

Map-Reduce for Data-Parallel ML

- Excellent for large data-parallel tasks!



Map Reduce

Feature
Extraction

Cross
Validation

Computing Sufficient
Statistics

Is there more to
Machine Learning

?

Concrete Example

Label Propagation

Label Propagation Algorithm

- Social Arithmetic:

50% What I list on my profile

40% Sue Ann Likes

+ 10% Carlos Like

I Like: 60% Cameras, 40% Biking

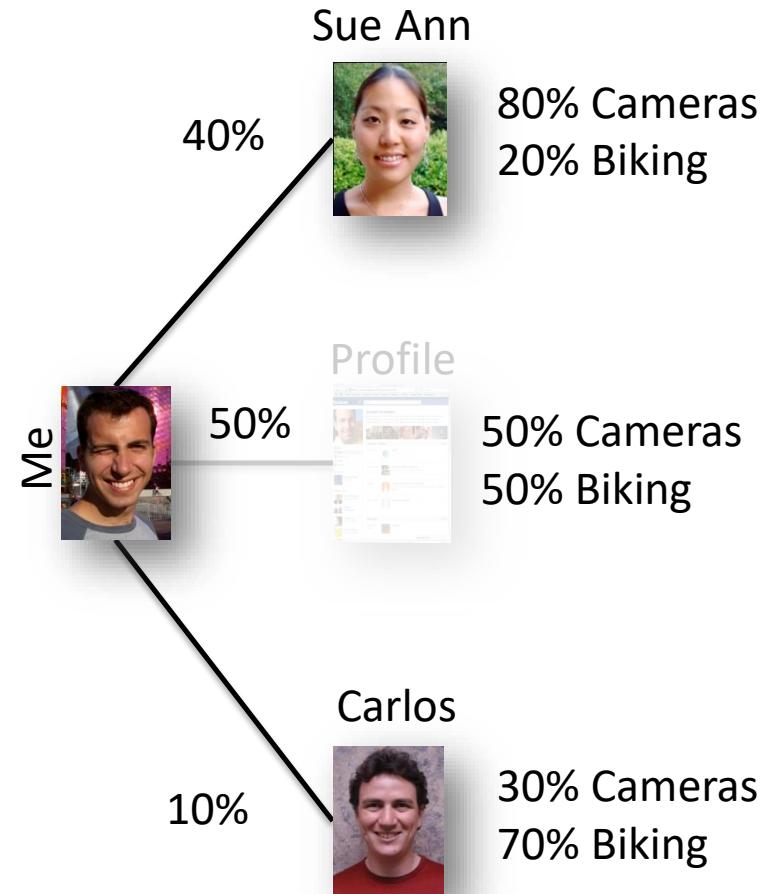
- Recurrence Algorithm:

$$Likes[i] = \sum_{j \in Friends[i]} W_{ij} \cdot Likes[j]$$

- iterate until convergence

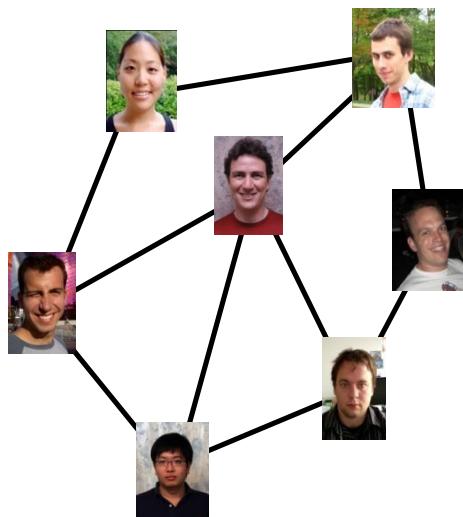
- Parallelism:

- Compute all $Likes[i]$ in parallel

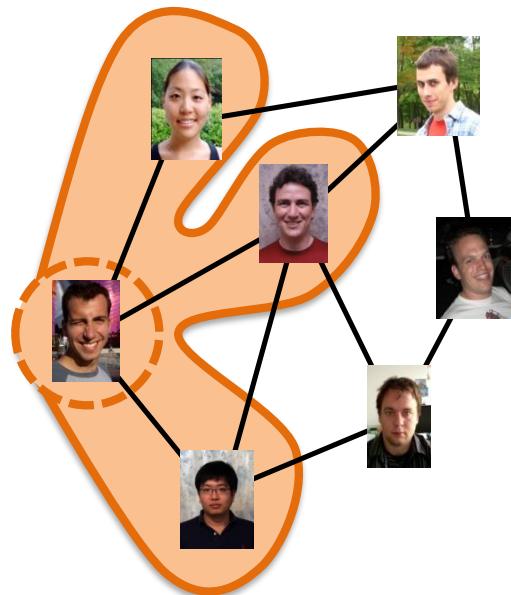


Properties of Graph Parallel Algorithms

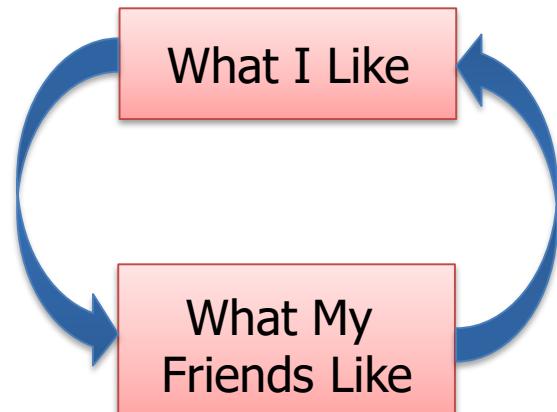
Dependency
Graph



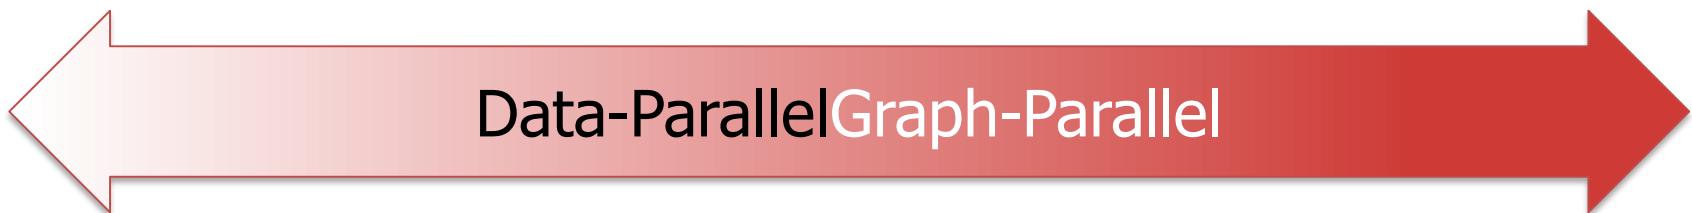
Factored
Computation



Iterative
Computation



- Excellent for large data-parallel tasks!



Map Reduce

Feature Extraction

Cross Validation

Computing Sufficient Statistics

Map Reduce?

Lasso

Tensor Factorization

Deep Belief Networks

Label Propagation

Kernel Methods

PageRank

Neural Networks

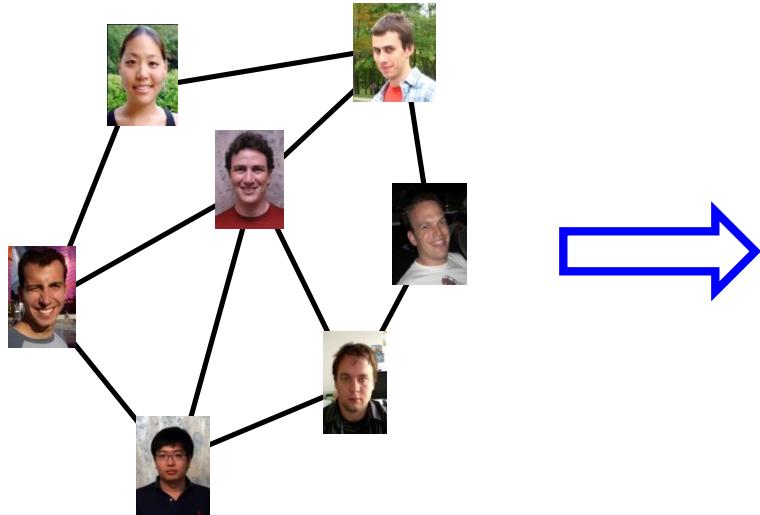
Belief

Propagation

*Why not use Map-Reduce
for
Graph Parallel Algorithms?*

Data Dependencies

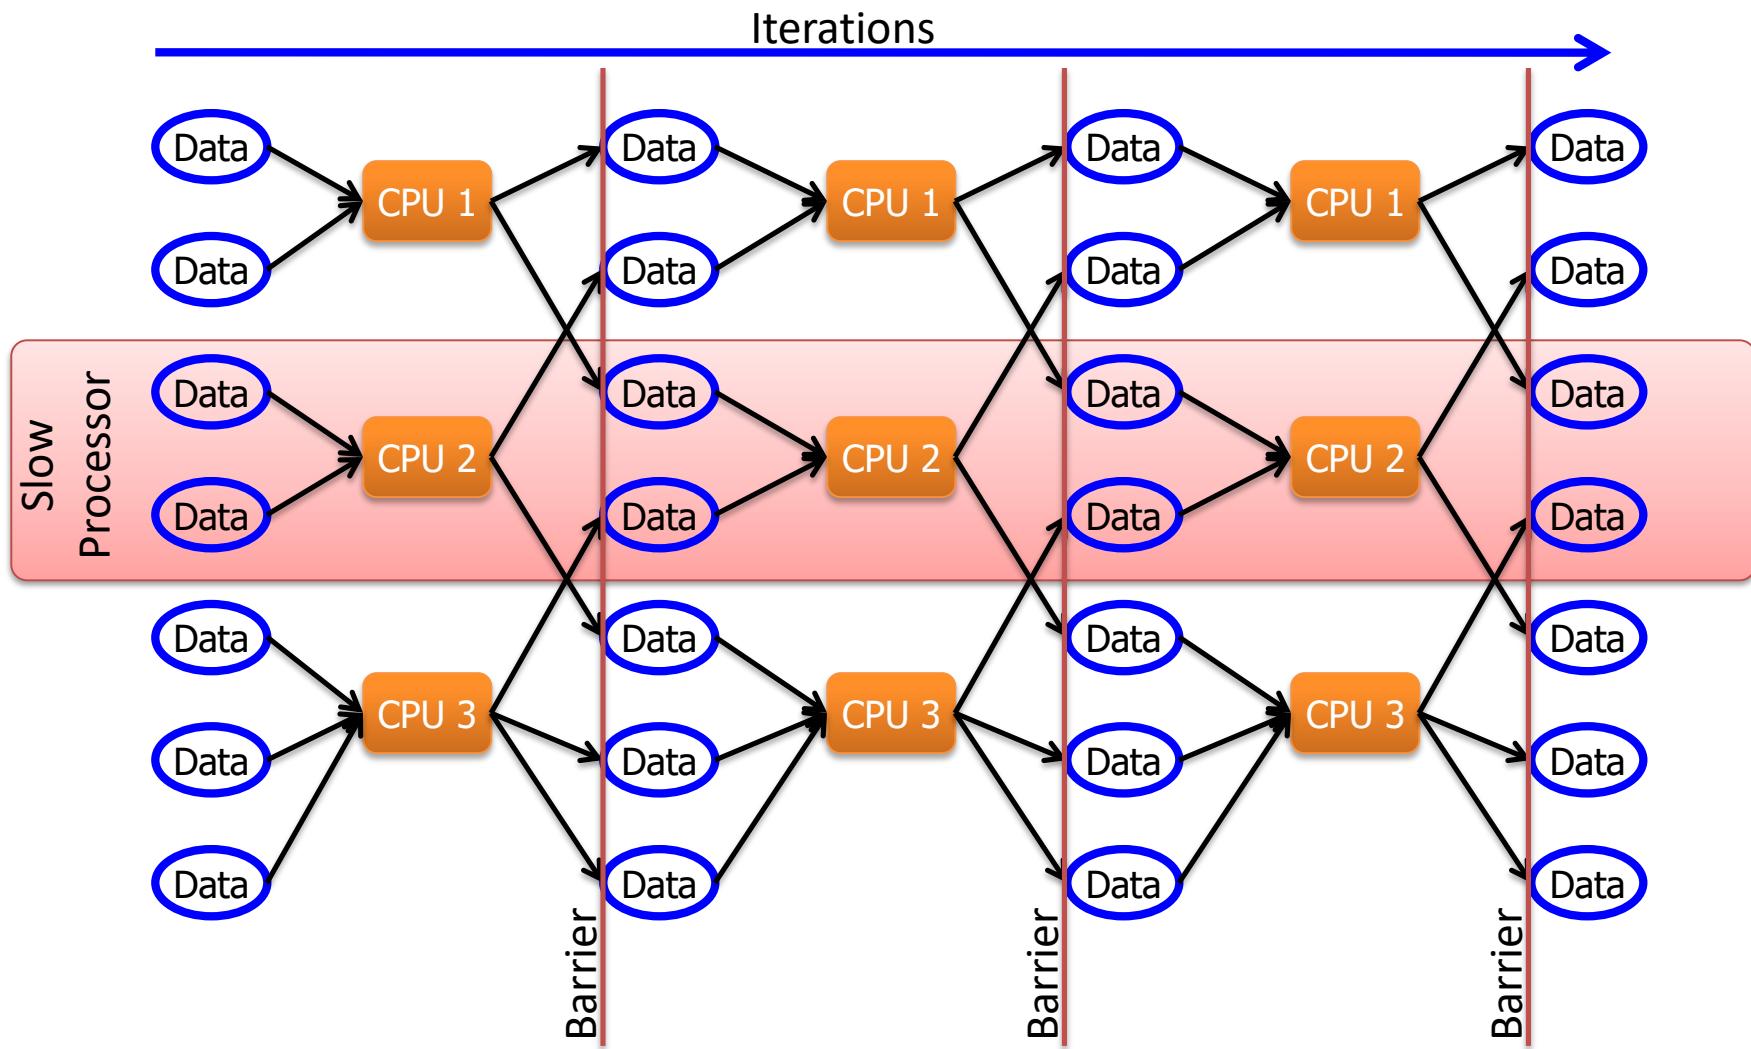
- Map-Reduce does not efficiently express dependent data
 - User must code substantial data transformations
 - Costly data replication



Independent Data Rows

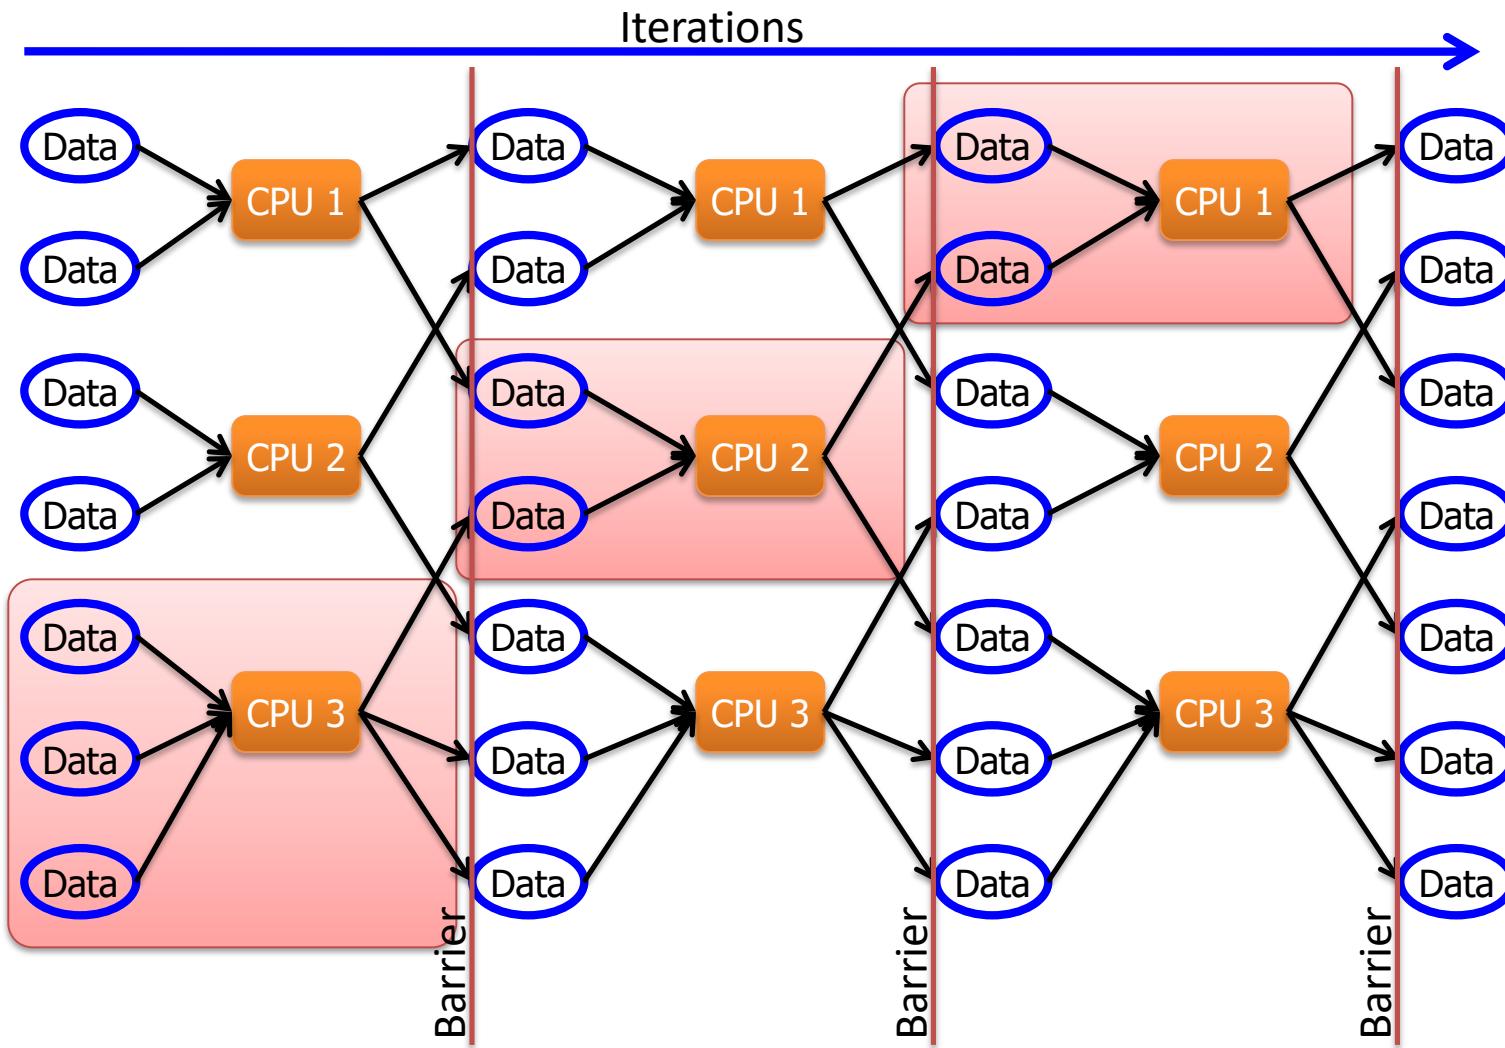
Iterative Algorithms

- Map-Reduce not efficiently express iterative algorithms:



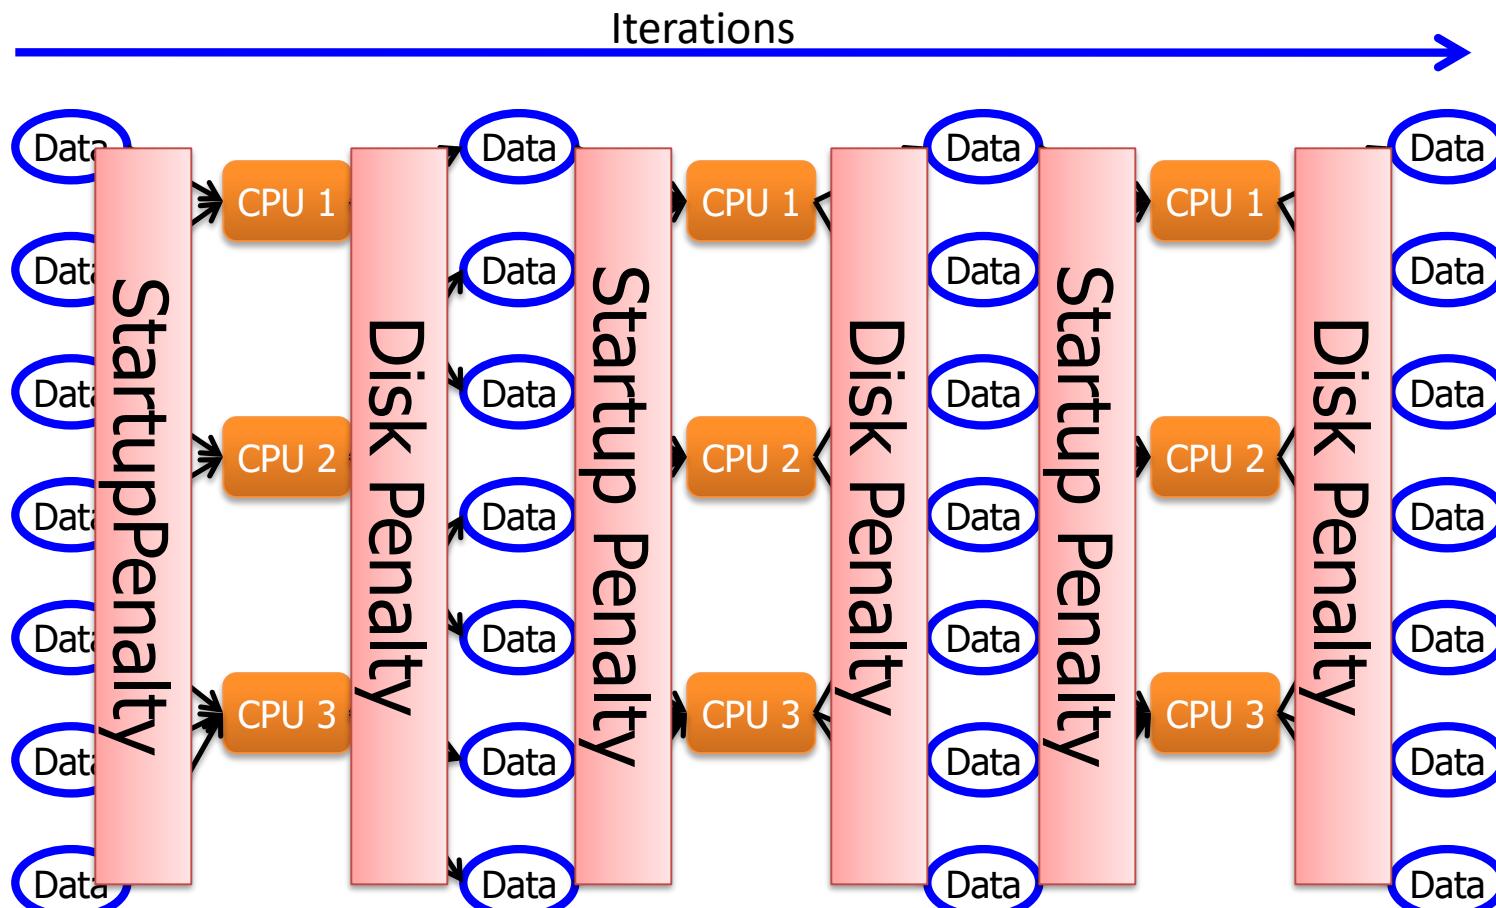
MapAbuse: Iterative MapReduce

- Only a subset of data needs computation:

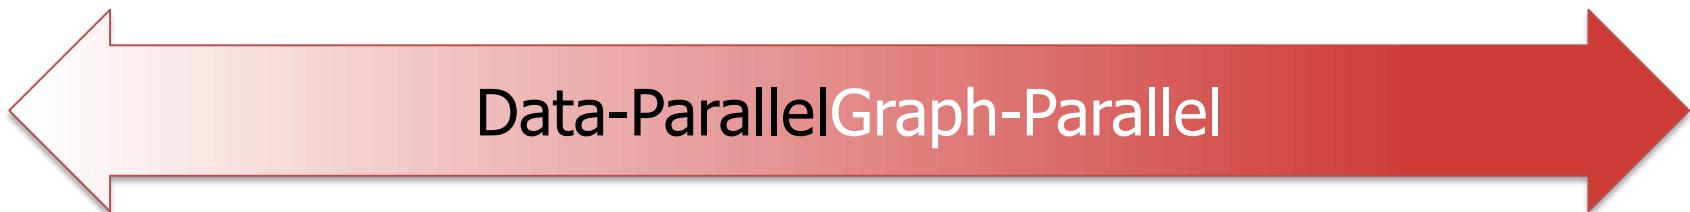


MapAbuse: Iterative MapReduce

- System is not optimized for iteration:



- Excellent for large data-parallel tasks!



Map Reduce

Feature Extraction

Computing Sufficient Statistics

Cross Validation

Pregel (Giraph)?

Lasso
Kernel Methods

Tensor Factorization

Deep Belief Networks

SVM

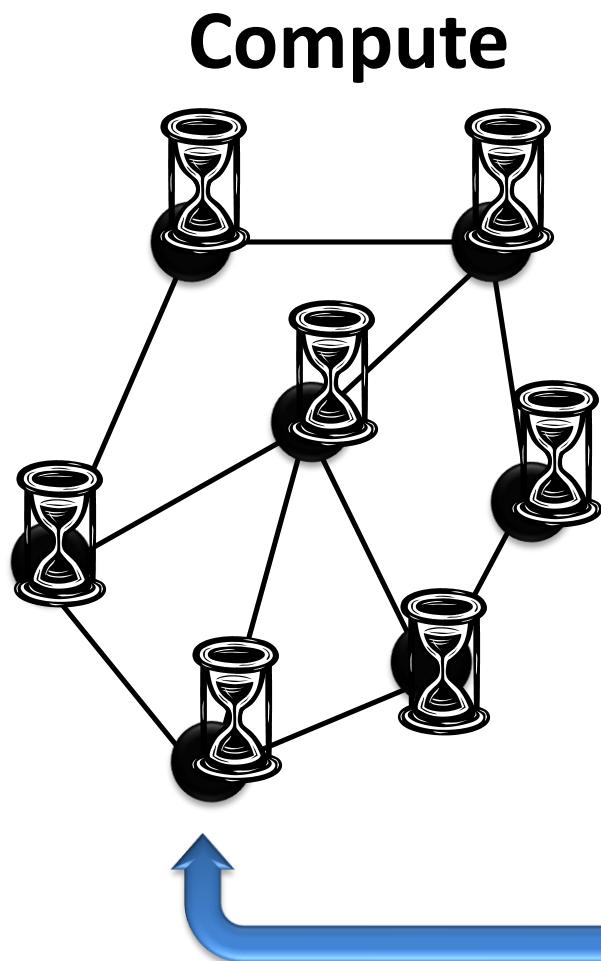
Belief Propagation

PageRank

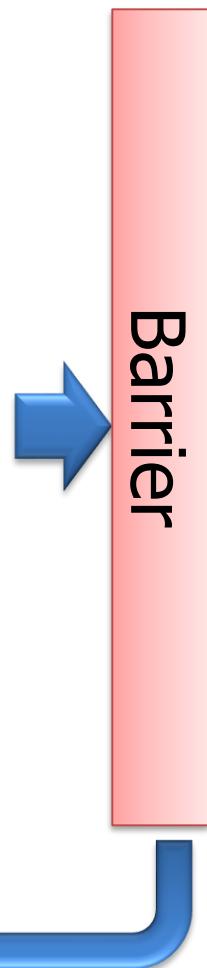
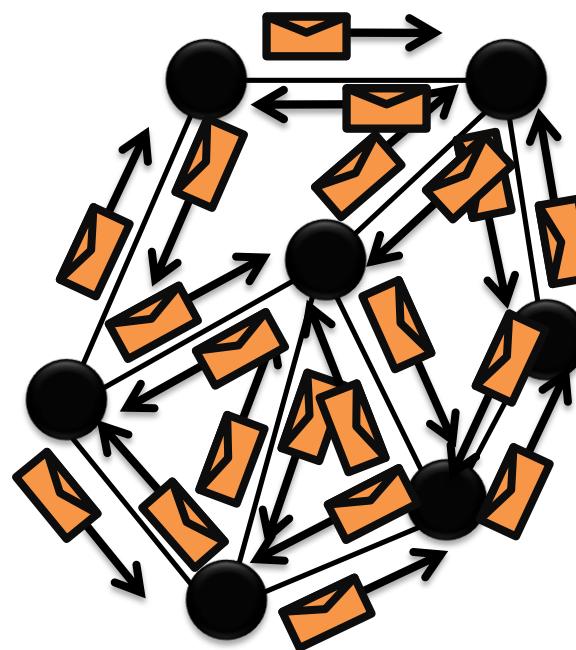
Neural Networks

Pregel (Giraph)

- Bulk Synchronous Parallel Model:



Communicate



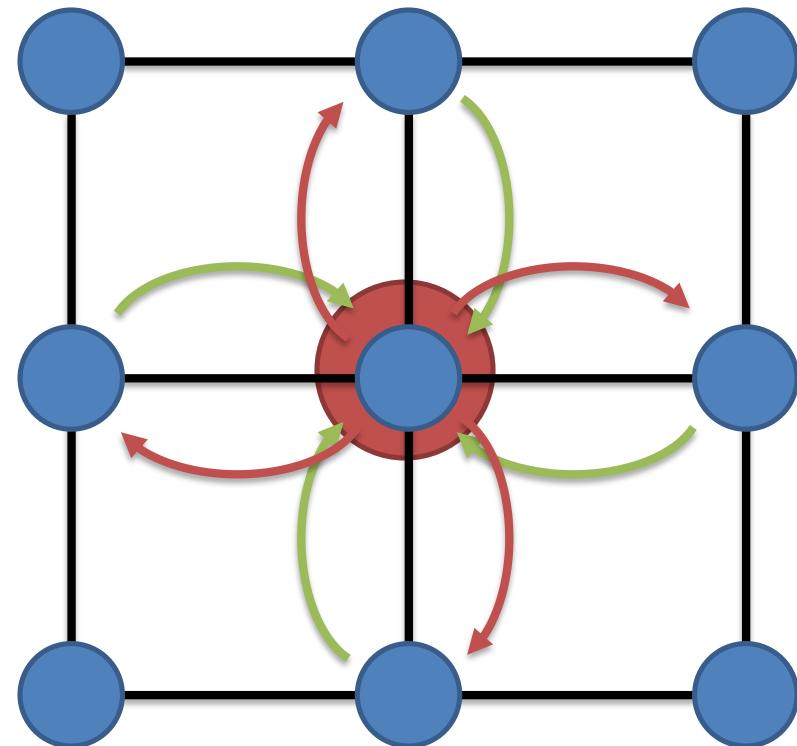
Problem

*Bulk synchronous computation
can be highly inefficient.*

Example:
Loopy Belief Propagation

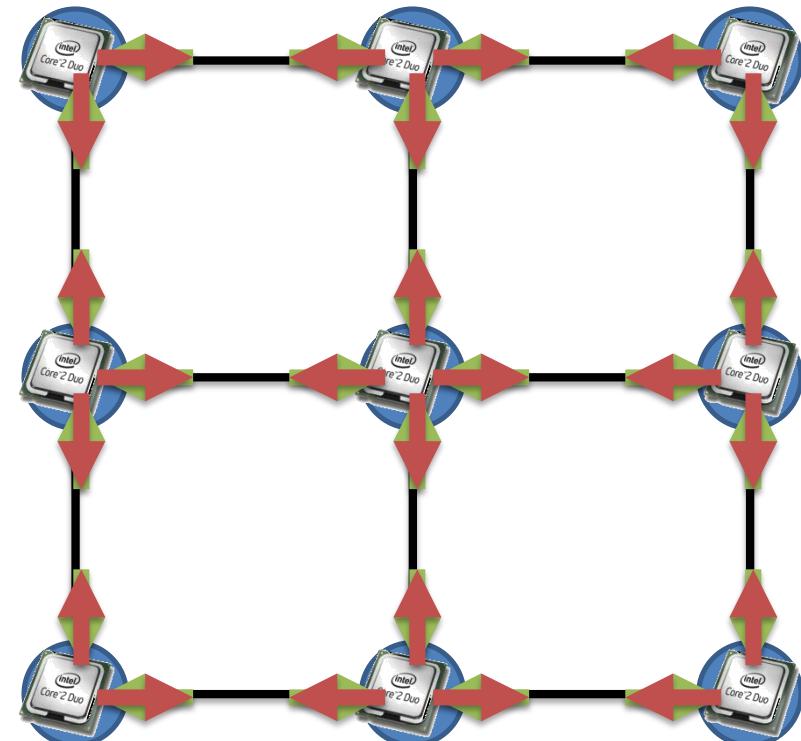
Loopy Belief Propagation (Loopy BP)

- Iteratively estimate the “beliefs” about vertices
 - Read **in messages**
 - Updates marginal estimate (**belief**)
 - Send updated **out messages**
- Repeat for all variables until convergence

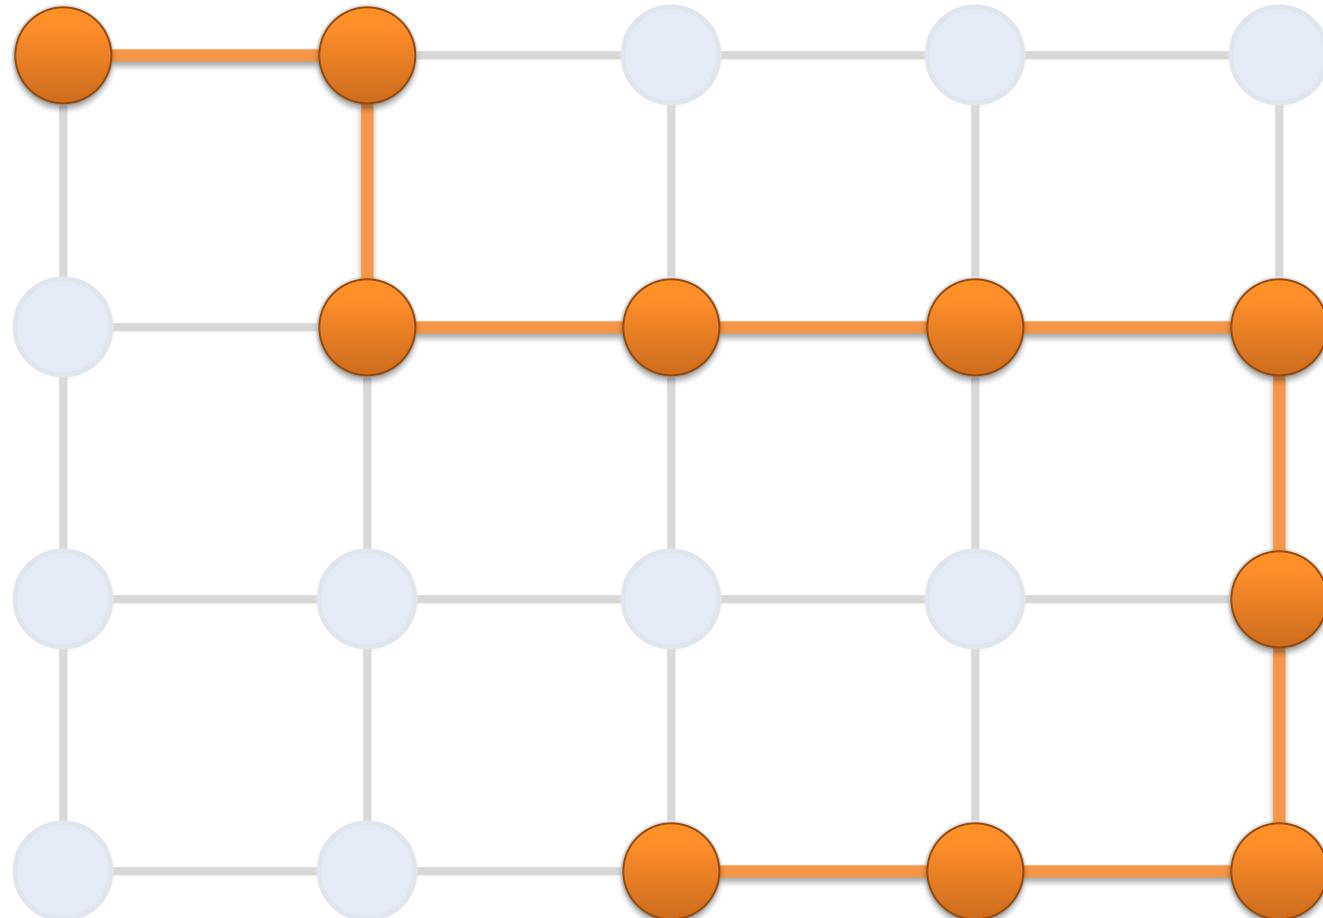


Bulk Synchronous Loopy BP

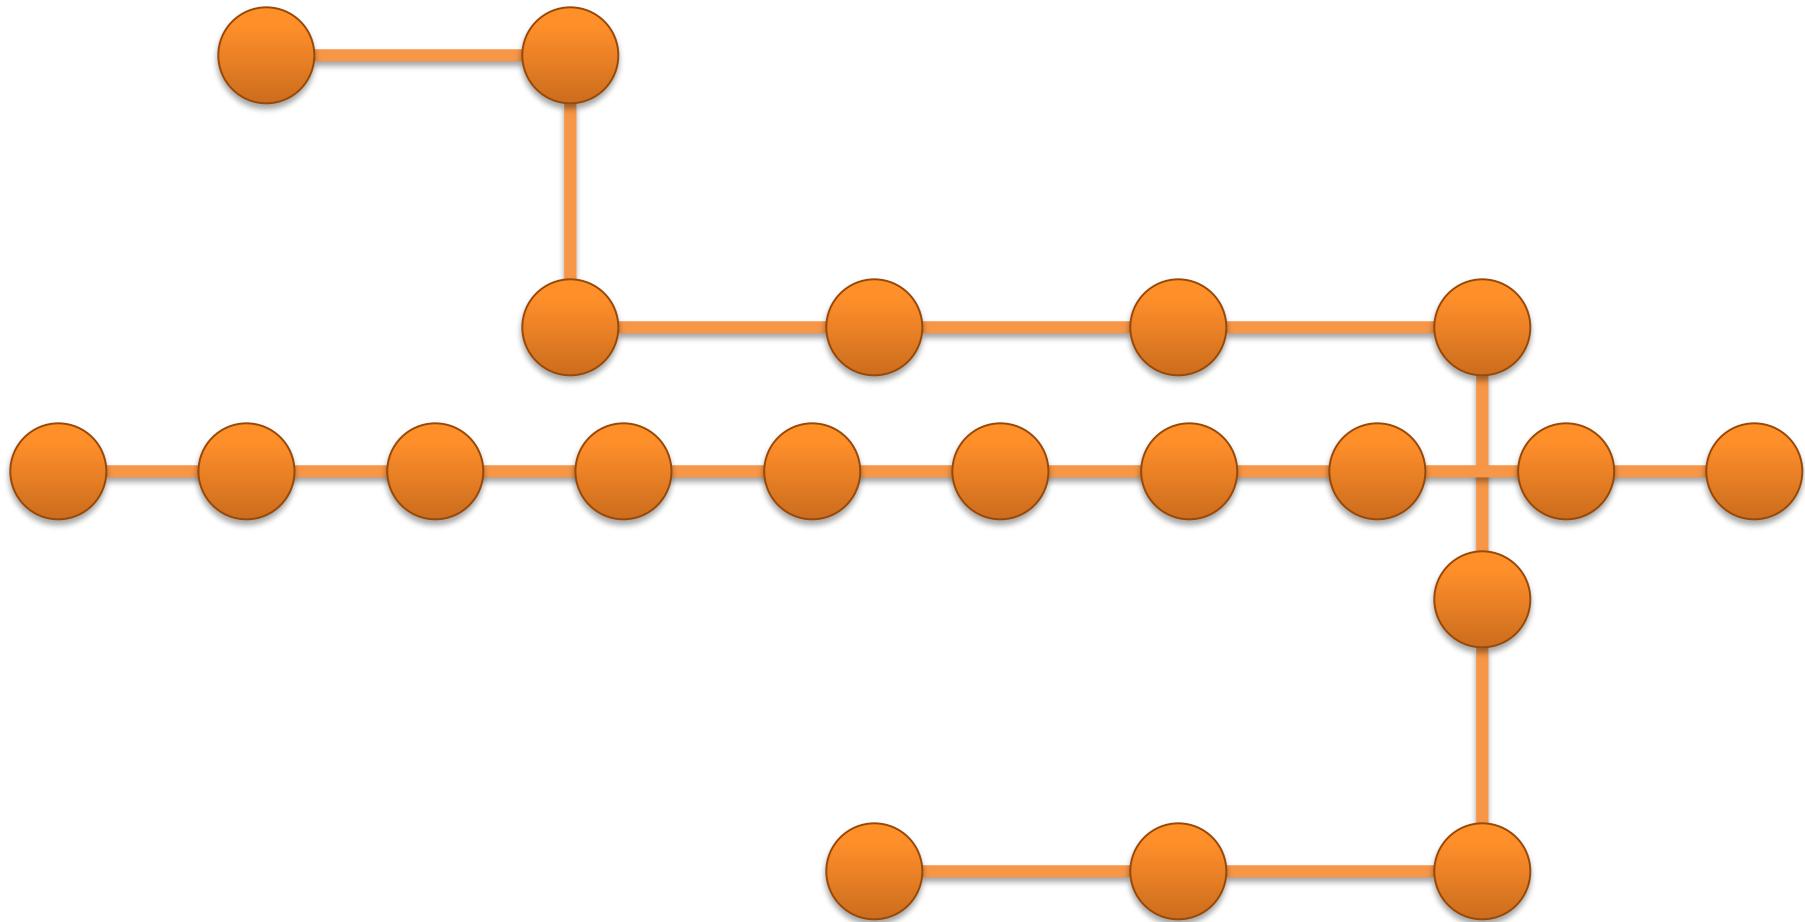
- Often considered embarrassingly parallel
 - Associate processor with each vertex
 - Receive all messages
 - Update all beliefs
 - Send all messages
- Proposed by:
 - Brunton et al. CRV'06
 - Mendiburu et al. GECC'07
 - Kang,et al. LDMTA'10
 - ...



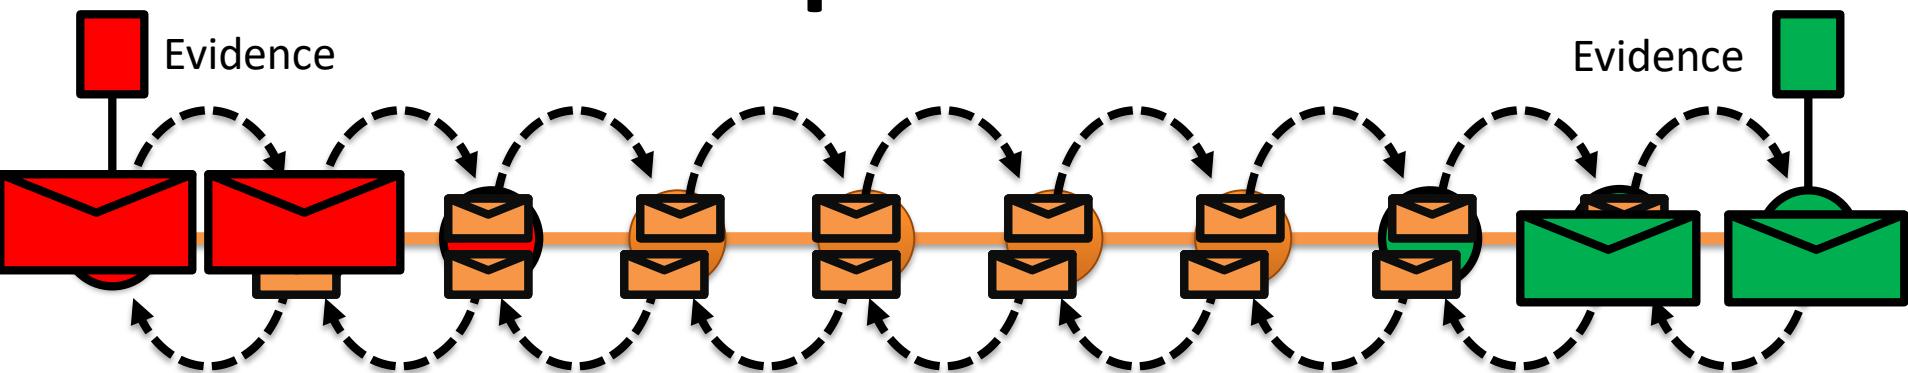
Sequential Computational Structure



Hidden Sequential Structure



Hidden Sequential Structure



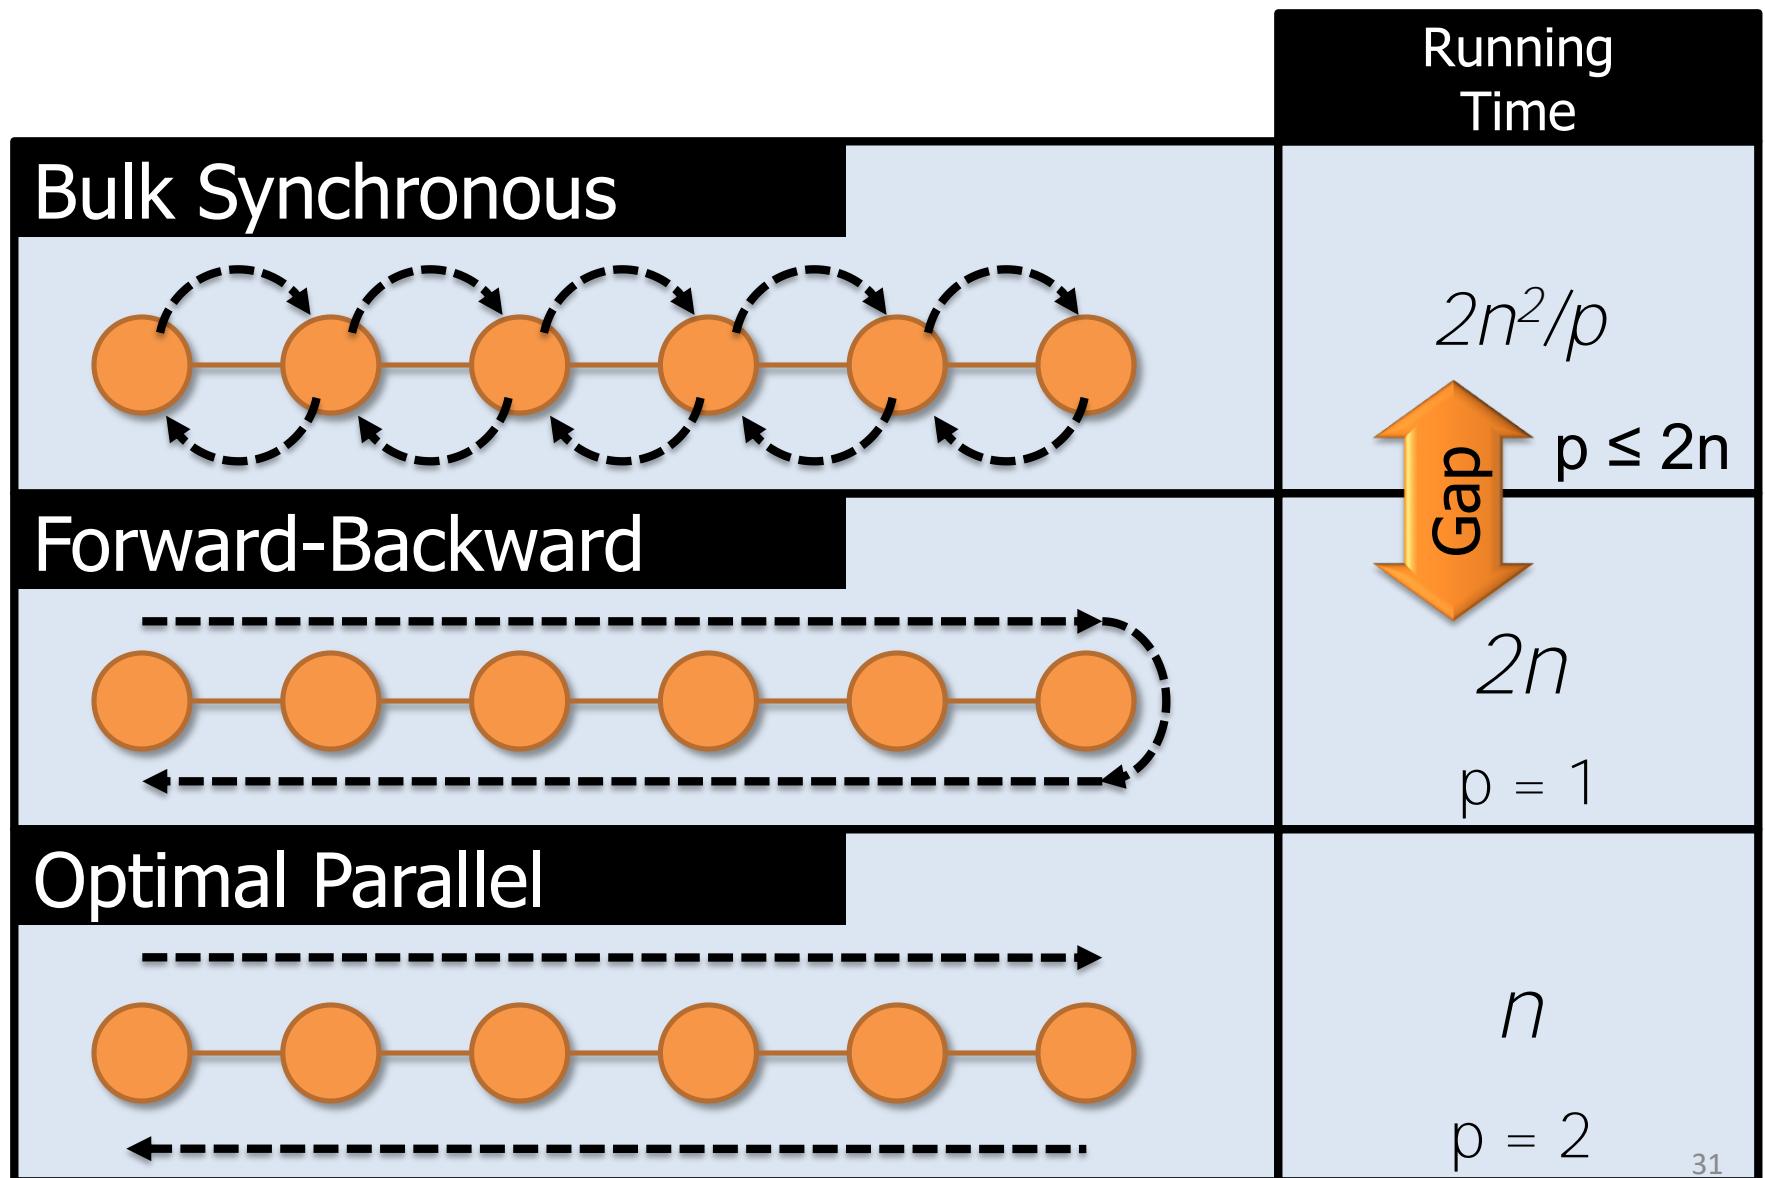
- Running Time:

$$\frac{2n \text{ Messages Calculations}}{p \text{ Processors}} \times (n \text{ Iterations to Converge}) = \frac{2n^2}{p}$$

Time for a single parallel iteration

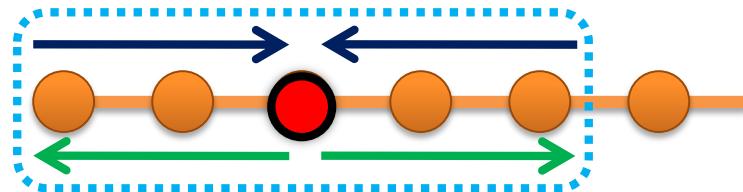
Number of Iterations

Optimal Sequential Algorithm



The Splash Operation

- Generalize the optimal chain algorithm:

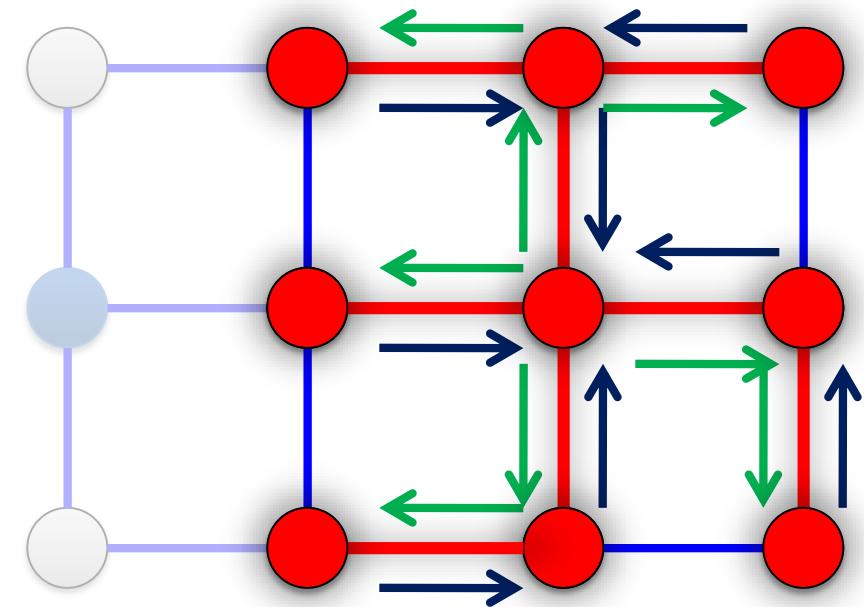


to arbitrary cyclic graphs:

1) Grow a BFS Spanning tree
with fixed size

2) Forward Pass computing all
messages at each vertex

3) Backward Pass computing all
messages at each vertex



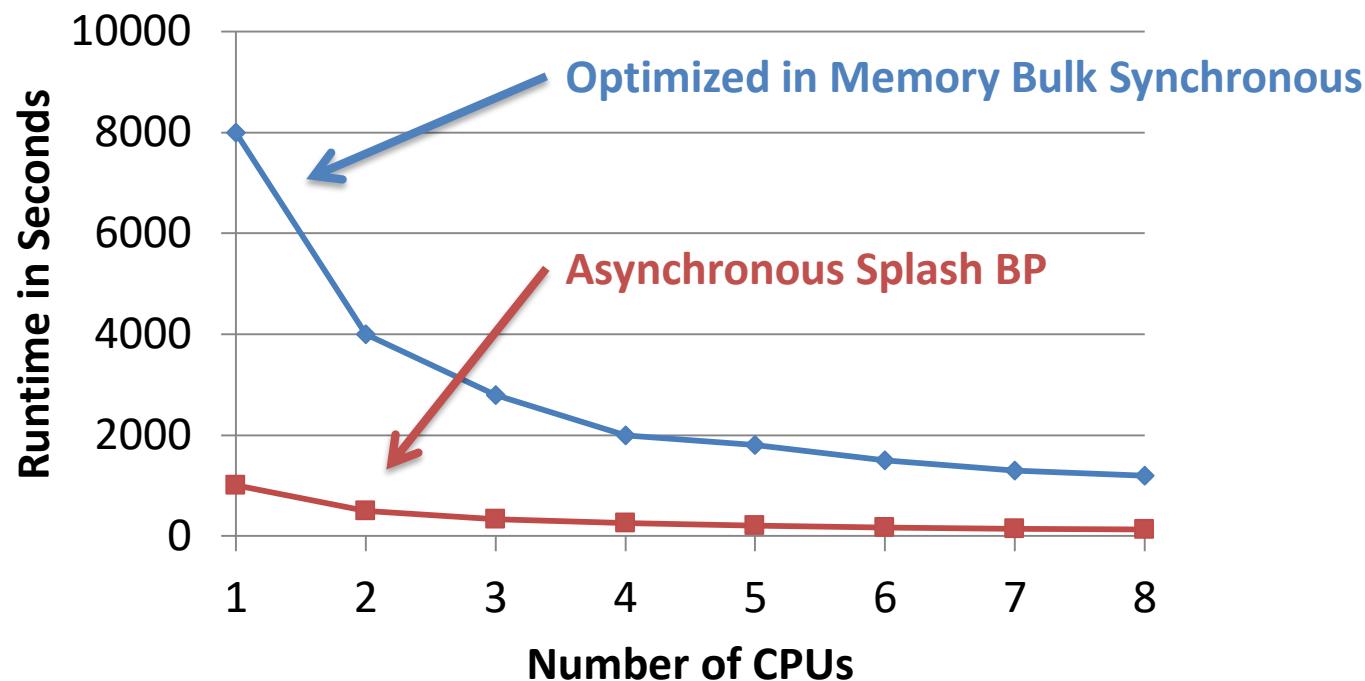
Data-Parallel Algorithms can be Inefficient

Residual Splash for Optimally Parallelizing Belief Propagation

Joseph E. Gonzalez
Carnegie Mellon University

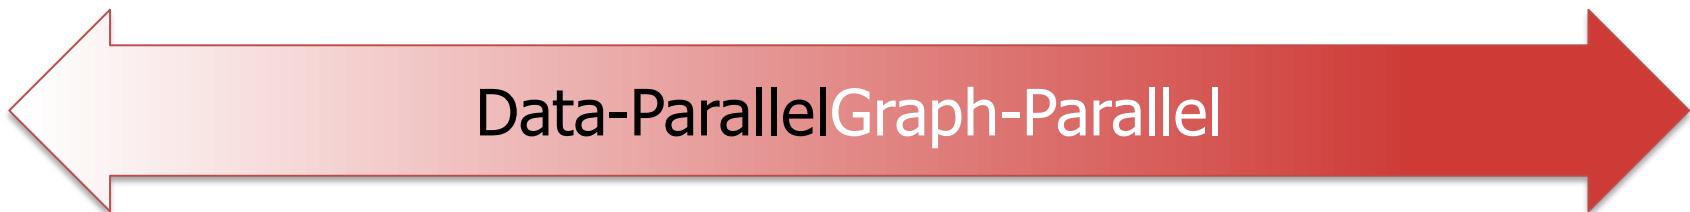
Yucheng Low
Carnegie Mellon University

Carlos Guestrin
Carnegie Mellon University



The limitations of the Map-Reduce abstraction can lead to inefficient parallel algorithms.

- Map-Reduce is not well suited for Graph-Parallelism



Map Reduce

Feature Extraction

Computing Sufficient Statistics

Cross Validation

Pregel (Giraph)?

Lasso
Kernel Methods

Tensor Factorization

Deep Belief Networks

SVM

Belief Propagation

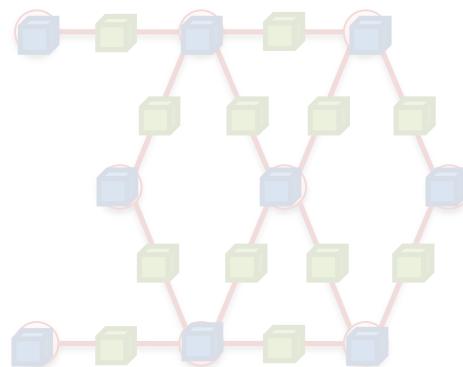
PageRank

Neural Networks

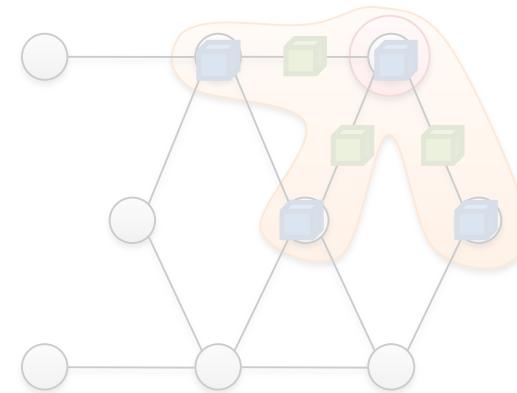
What is GraphLab?

The GraphLab Framework

Graph Based
Data Representation



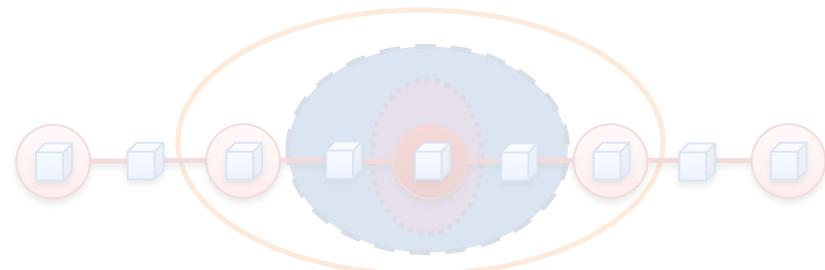
Update Functions
User Computation



Scheduler

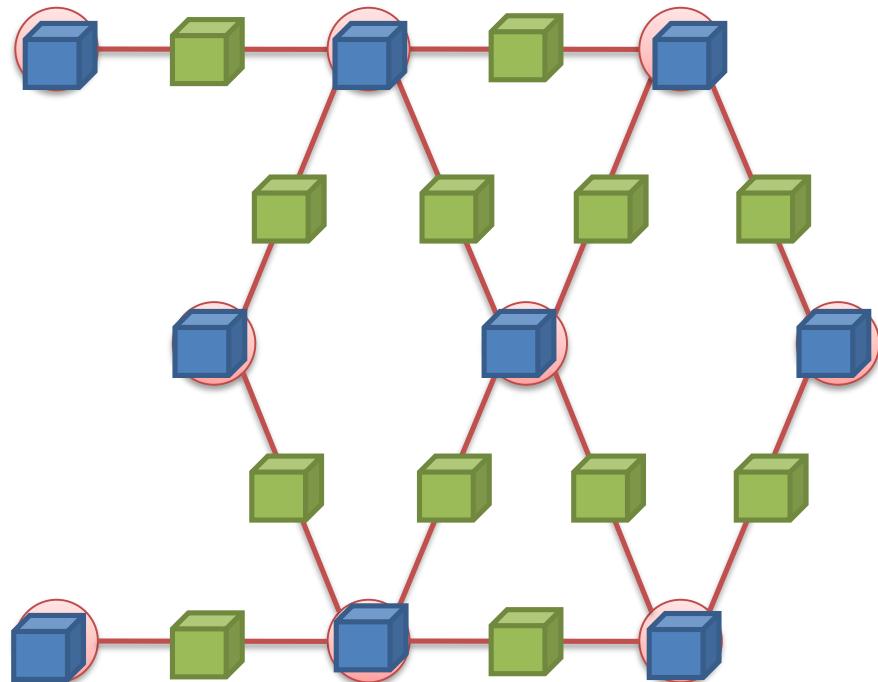


Consistency Model



Data Graph

A **graph** with arbitrary data (C++ Objects) associated with each vertex and edge.



Graph:

- Social Network

Vertex Data:

- User profile text
- Current interests estimates

Edge Data:

- Similarity weights

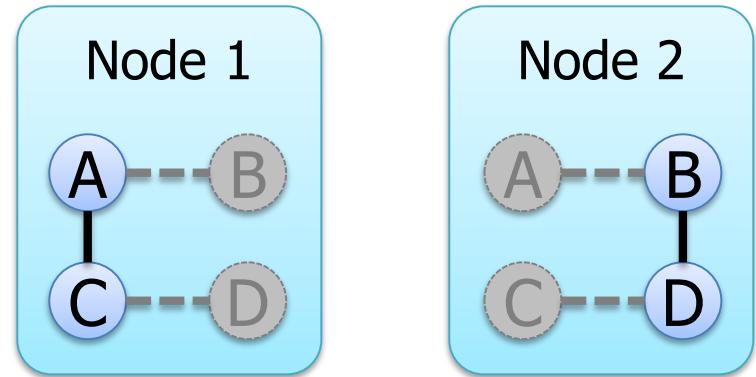
Implementing the Data Graph

Multicore Setting

- In Memory
- Relatively Straight Forward
 - `vertex_data(vid) → data`
 - `edge_data(vid,vid) → data`
 - `neighbors(vid) → vid_list`
- Challenge:
 - Fast lookup, low overhead
- Solution:
 - Dense data-structures
 - Fixed Vdata&Edata types
 - Immutable graph structure

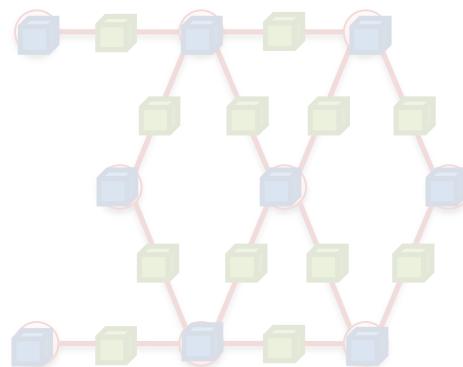
Cluster Setting

- In Memory
- Partition Graph:
 - ParMETIS or Random Cuts
- Cached Ghosting

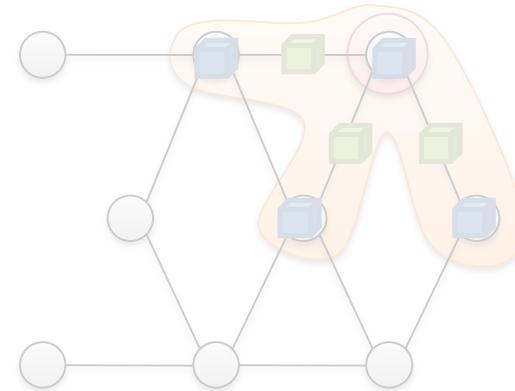


The GraphLab Framework

Graph Based
Data Representation



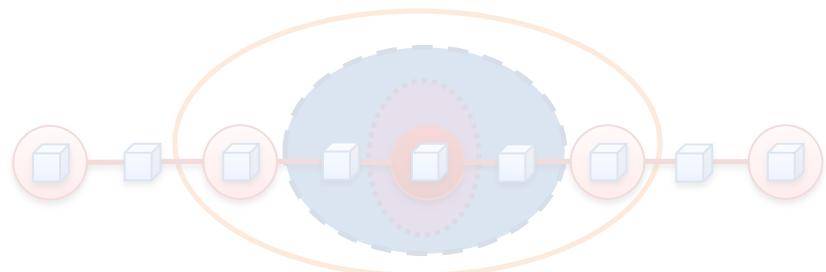
Update Functions
User Computation



Scheduler

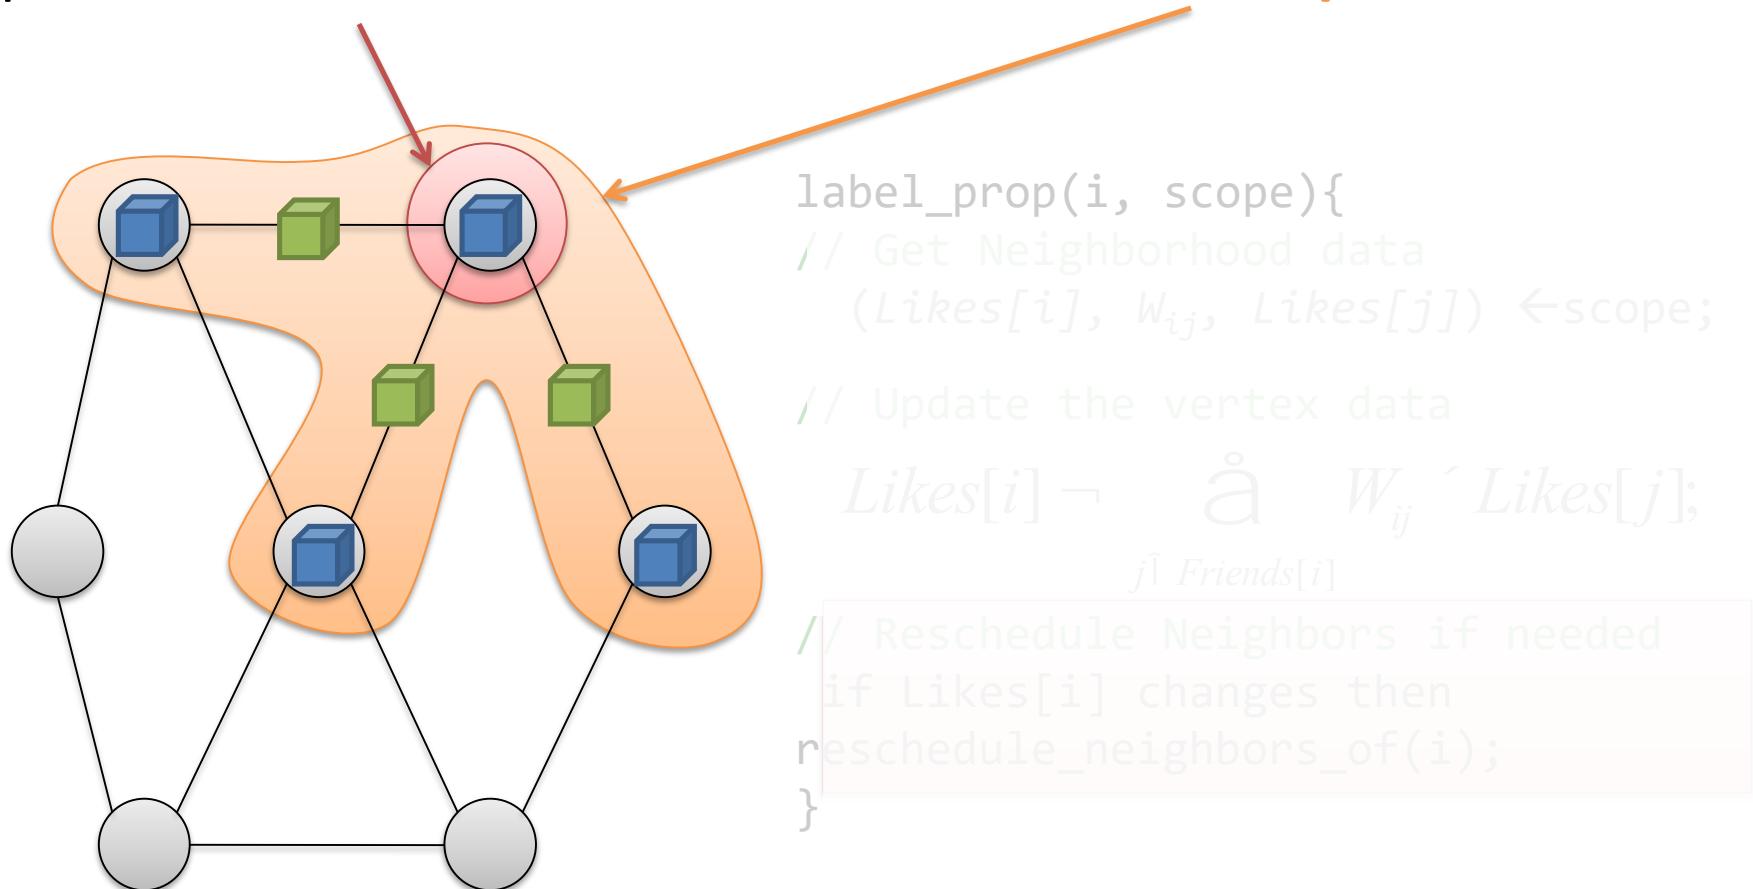


Consistency Model



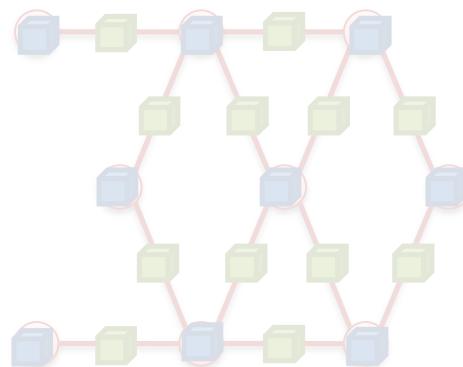
Update Functions

An **update function** is a user defined program which when applied to a **vertex** transforms the data in the **scope** of the vertex

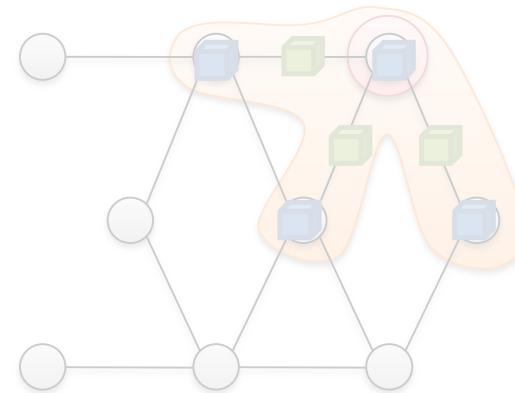


The GraphLab Framework

Graph Based
Data Representation



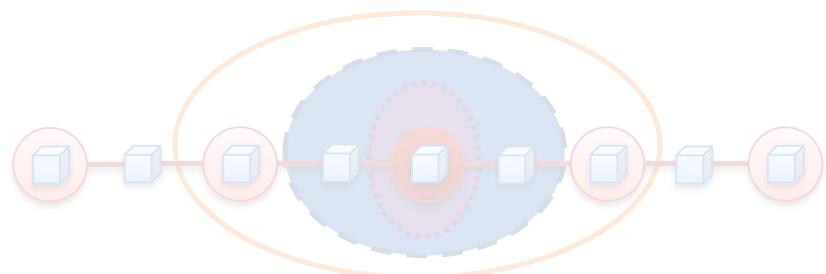
Update Functions
User Computation



Scheduler

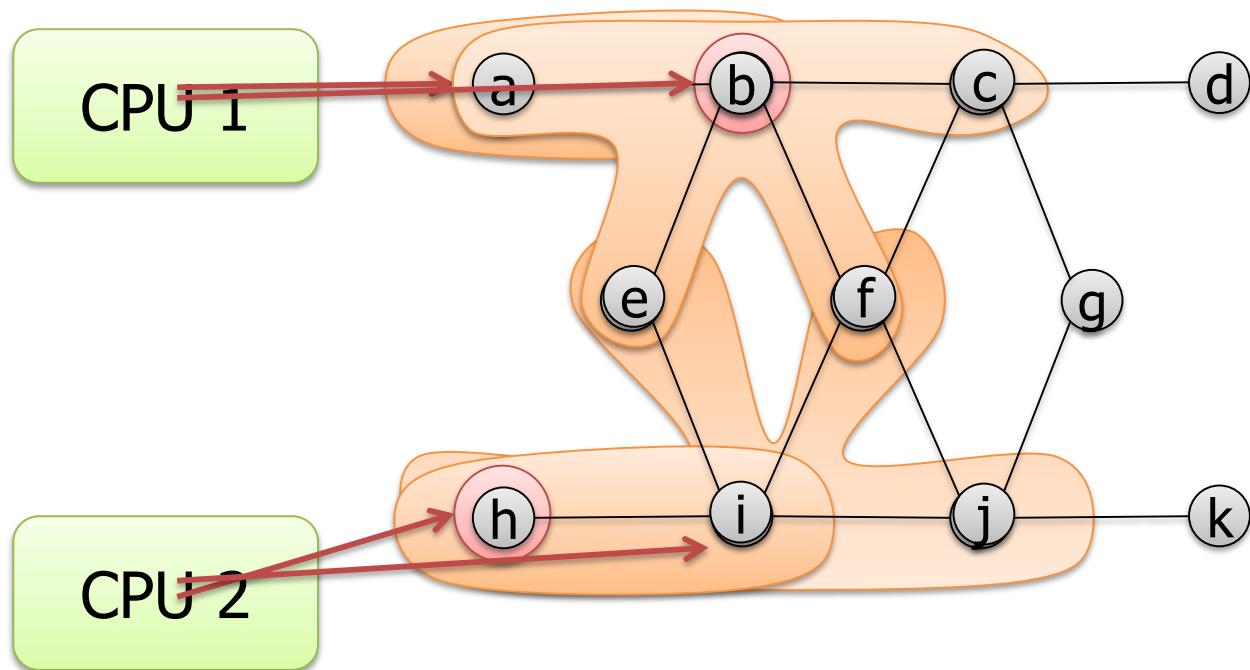
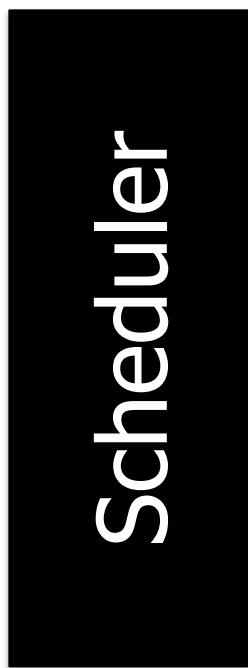


Consistency Model



The Scheduler

The **scheduler** determines the order that vertices are updated.

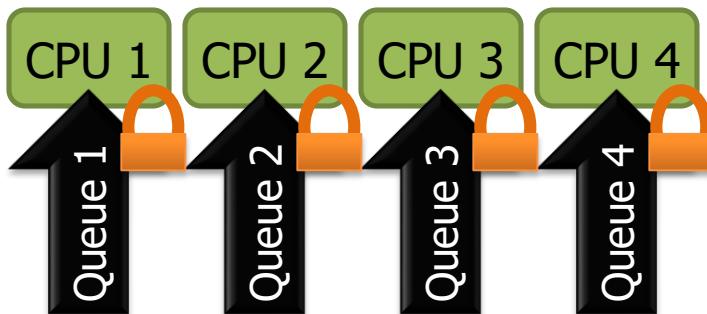


The process repeats until the scheduler is empty.

Implementing the Schedulers

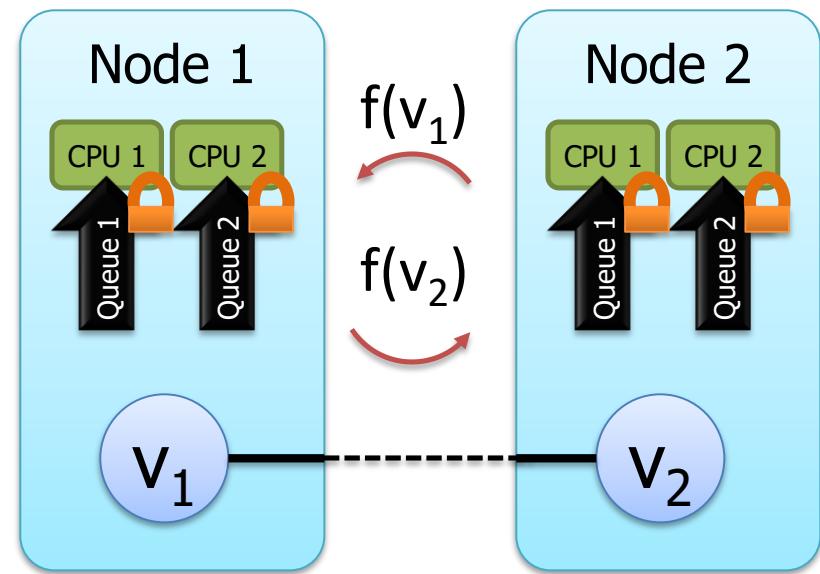
Multicore Setting

- Challenging!
 - Fine-grained locking
 - Atomic operations
- Approximate FiFo/Priority
 - Random placement
 - Work stealing



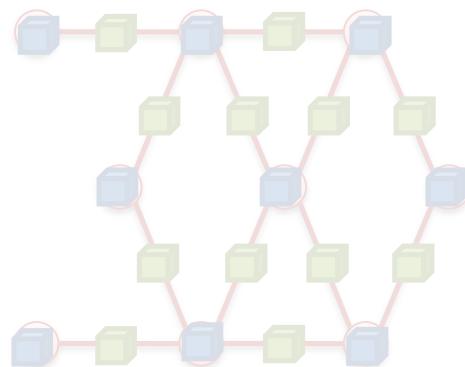
Cluster Setting

- Multicore scheduler on each node
 - Schedules only “local” vertices
 - Exchange update functions

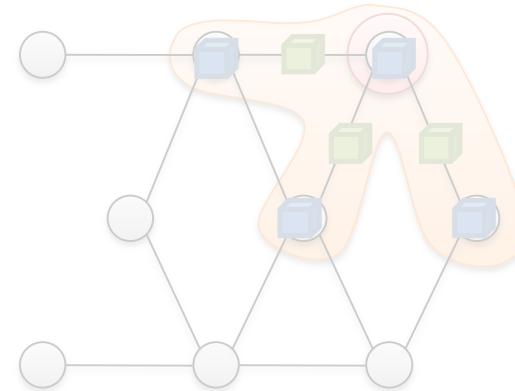


The GraphLab Framework

Graph Based
Data Representation



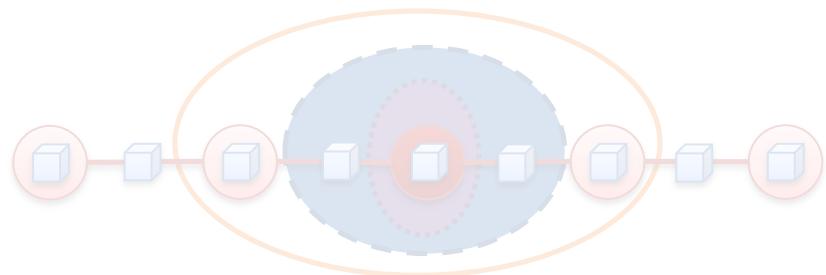
Update Functions
User Computation



Scheduler

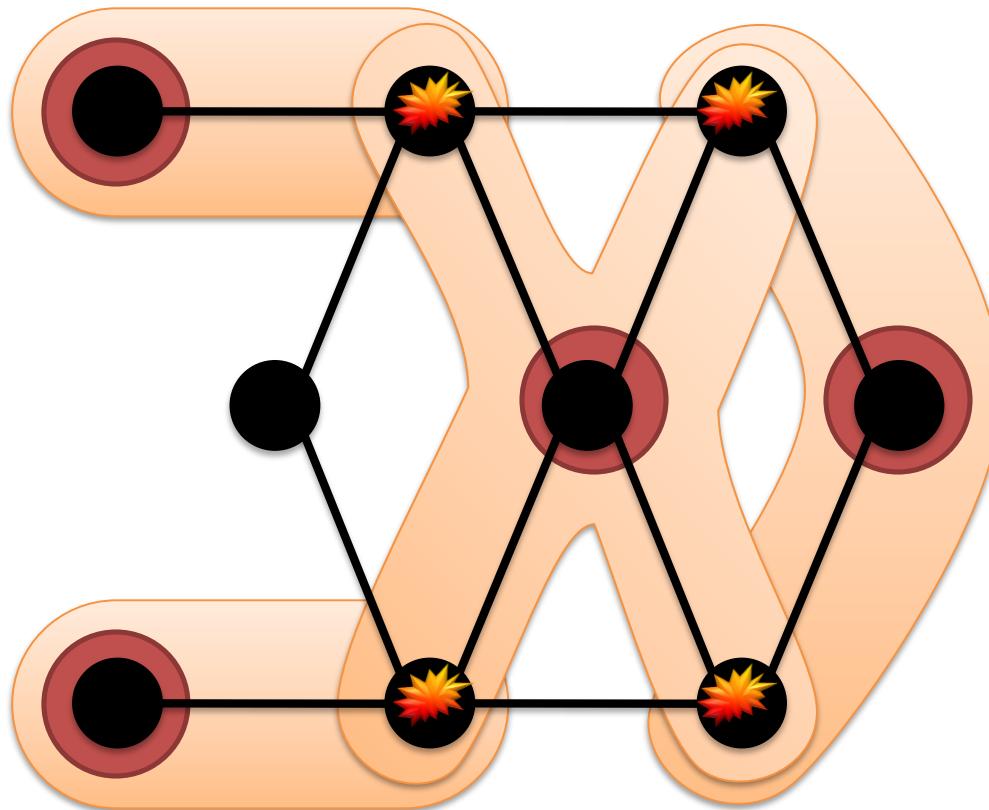


Consistency Model



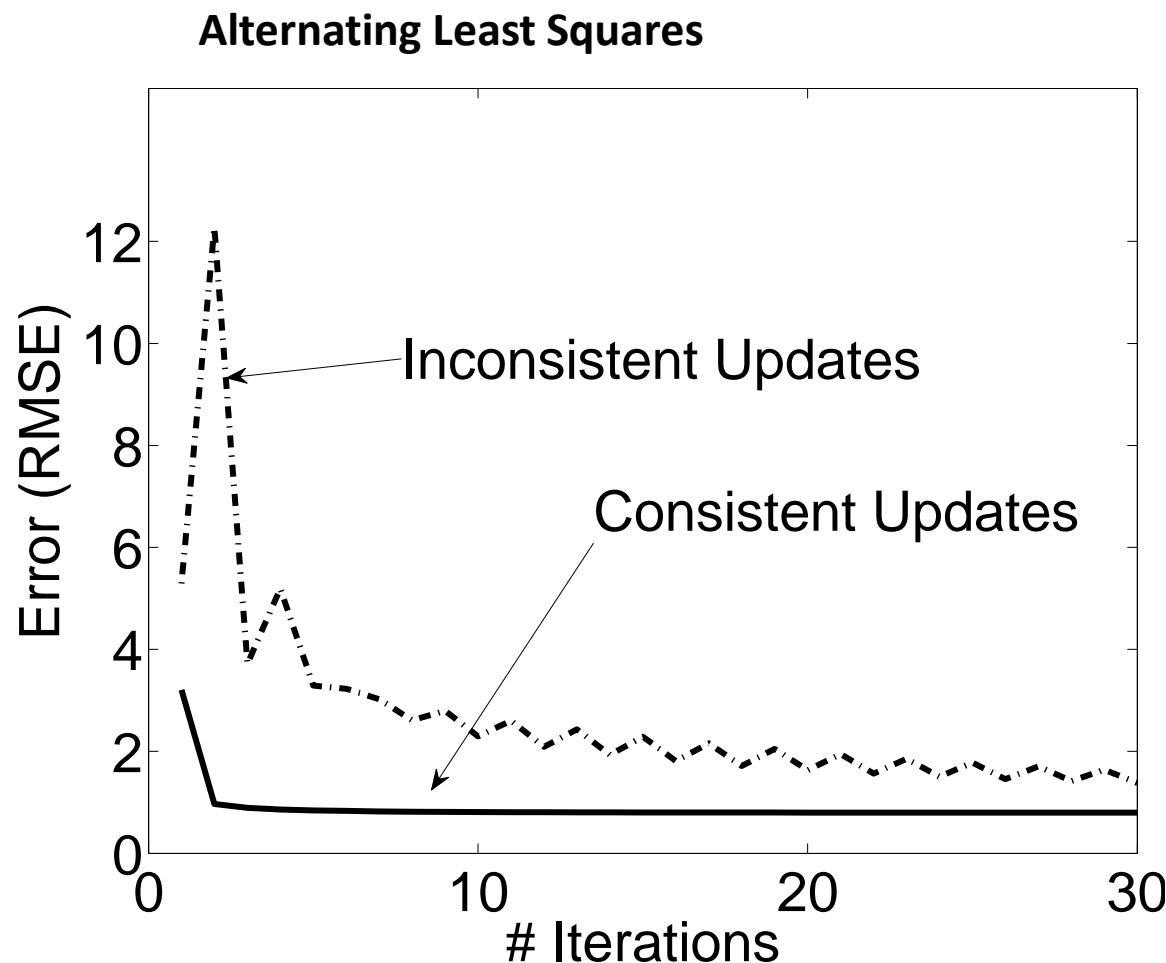
Ensuring Race-Free Code

- How much can computation **overlap**?



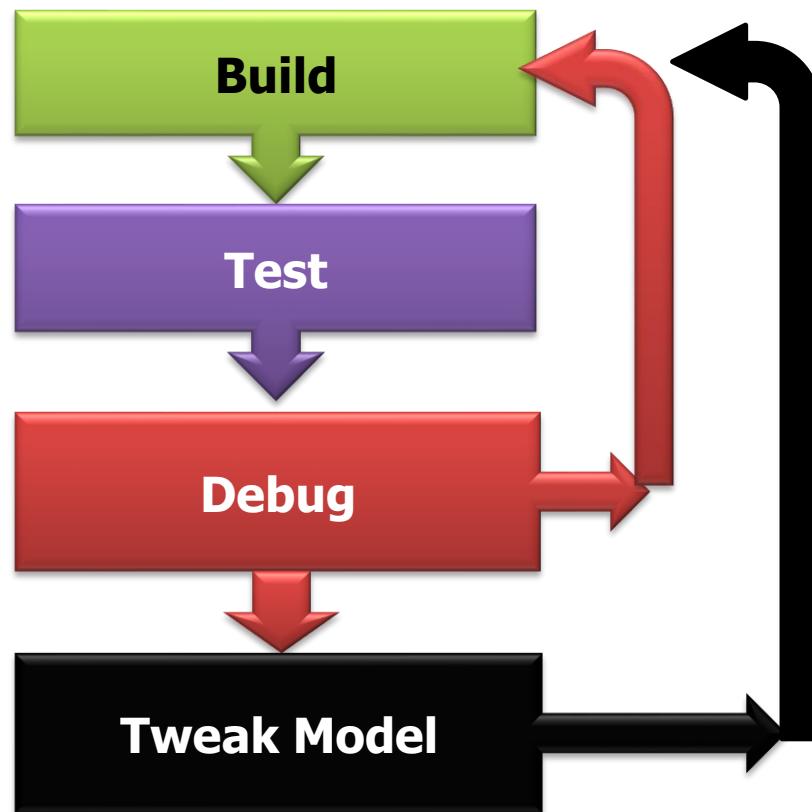
Importance of Consistency

Many algorithms require strict consistency, or performs significantly better under strict consistency.



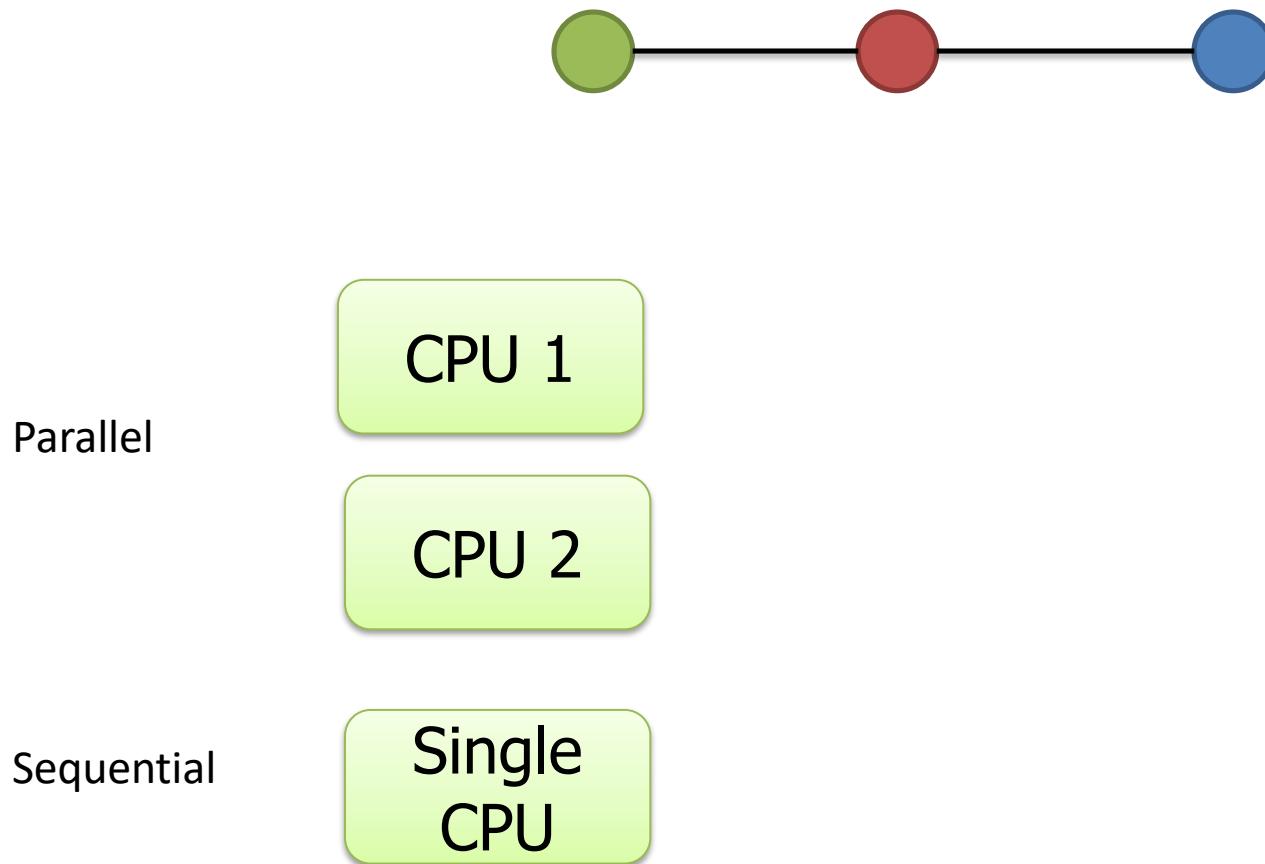
Importance of Consistency

Machine learning algorithms require “model debugging”

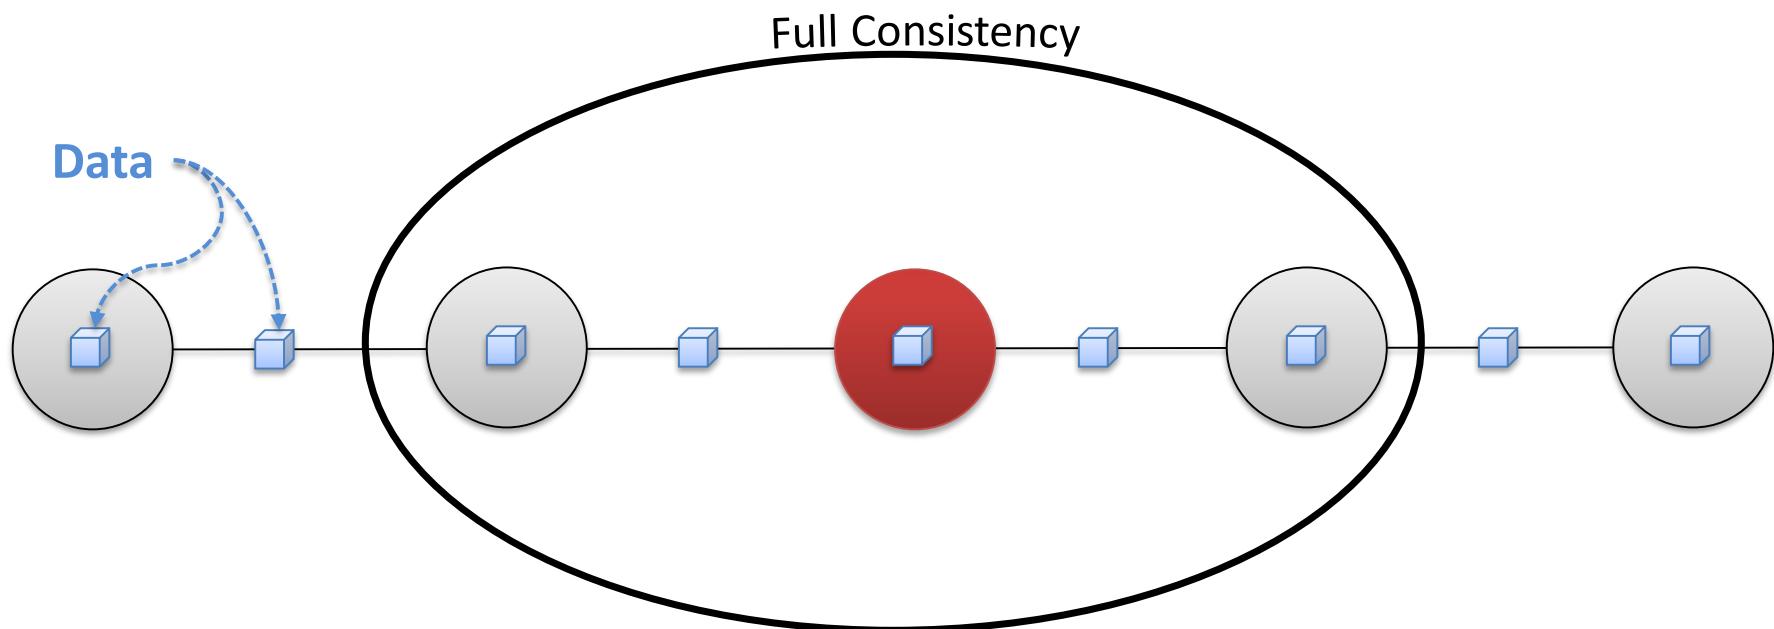


GraphLab Ensures Sequential Consistency

For each **parallel execution**, there exists a **sequential execution** of update functions which produces the same result.

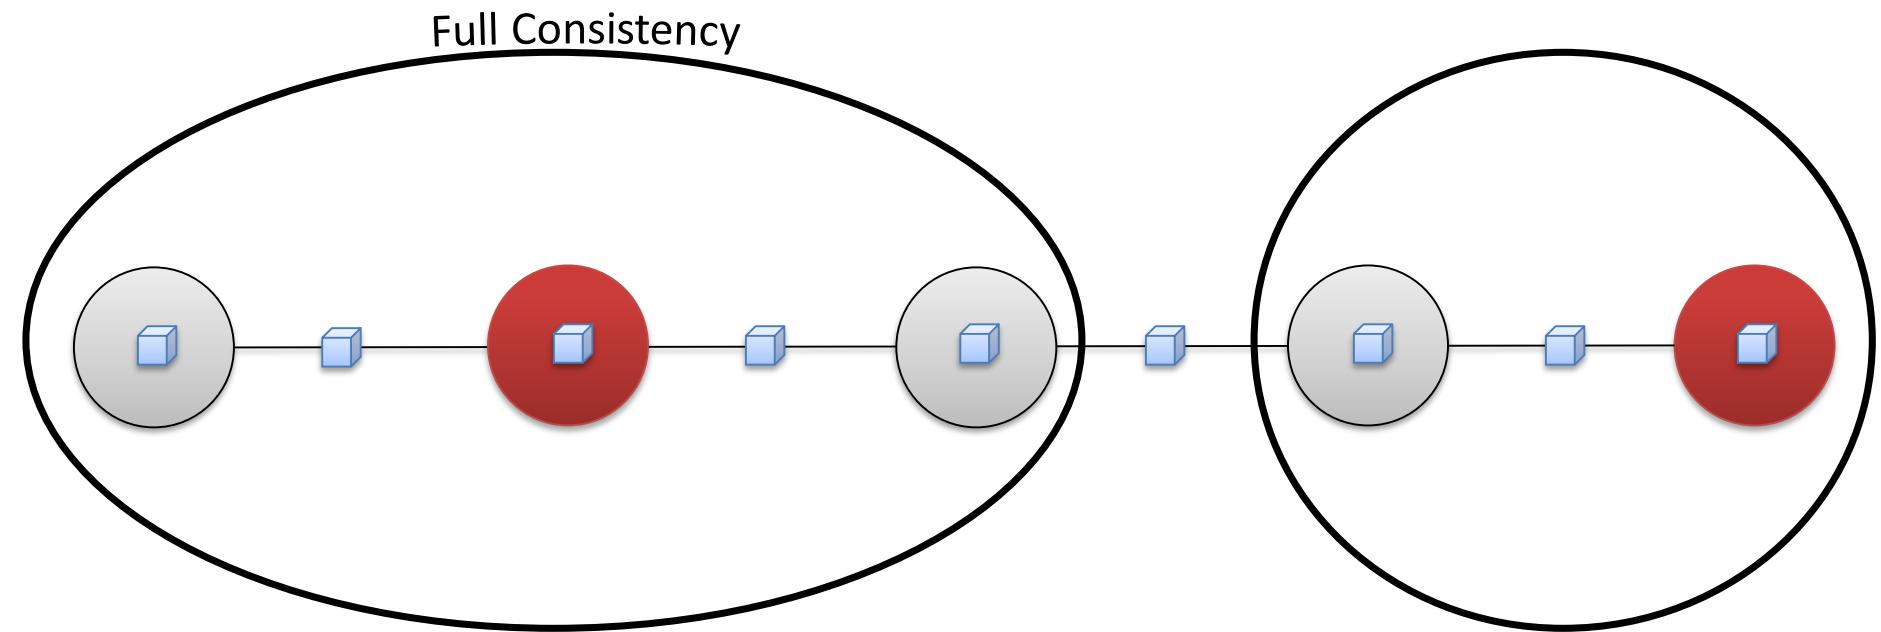


Consistency Rules

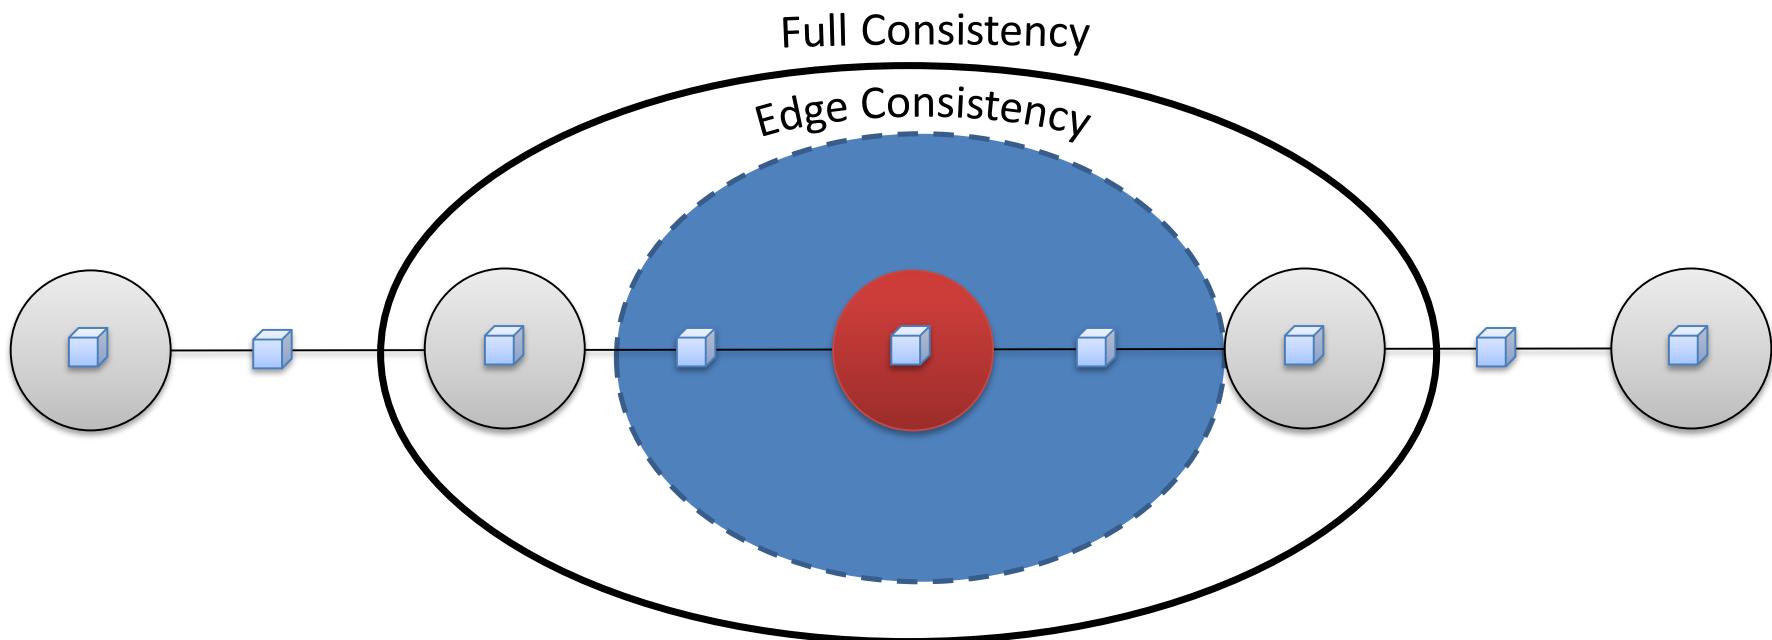


Guaranteed sequential consistency for all update functions

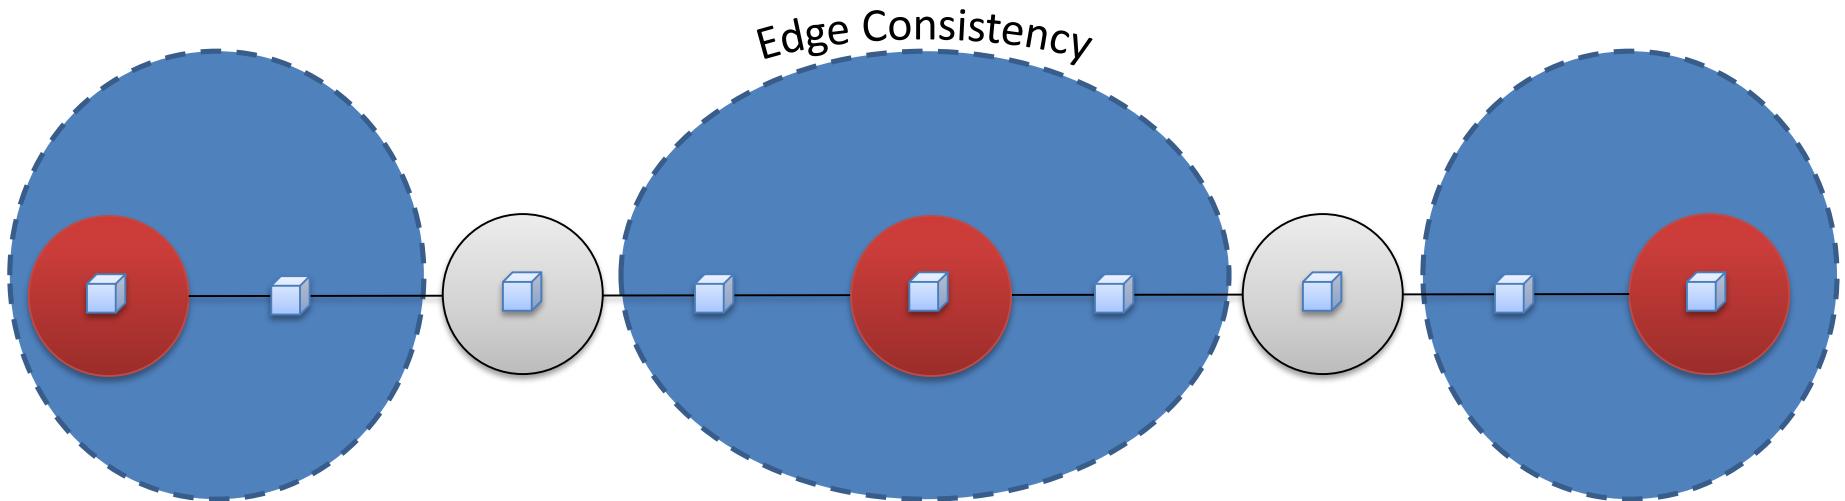
Full Consistency



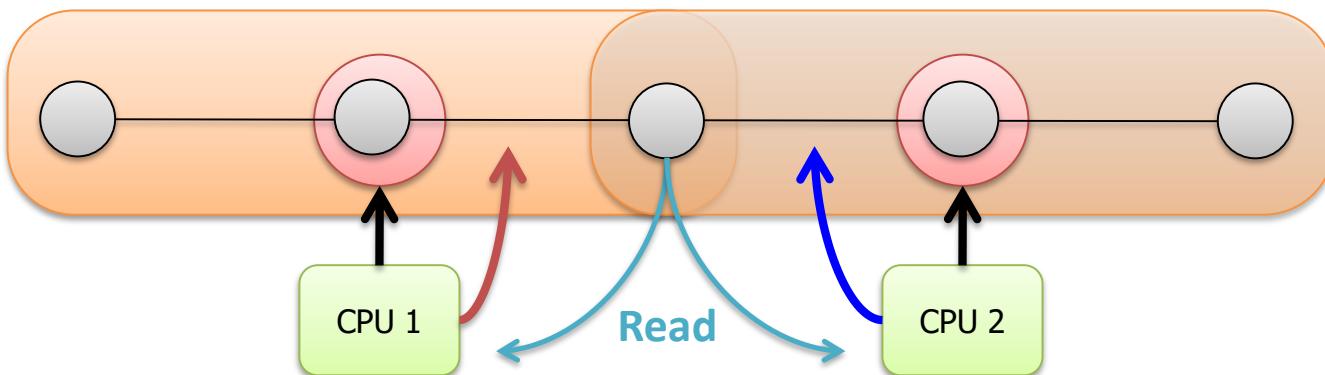
Obtaining More Parallelism

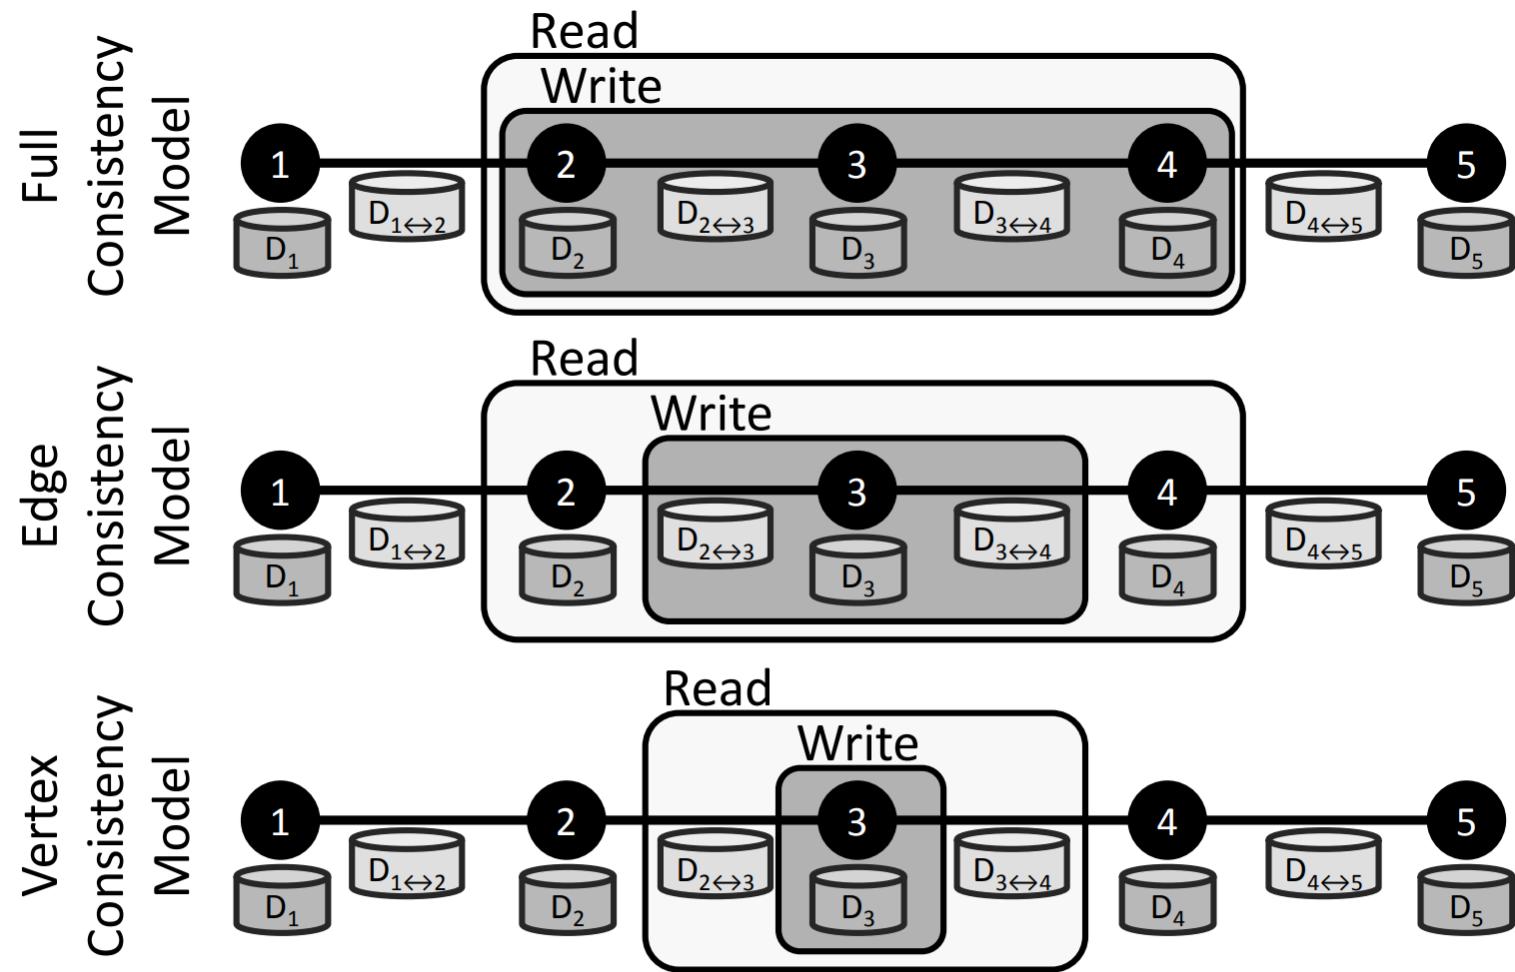


Edge Consistency



Safe





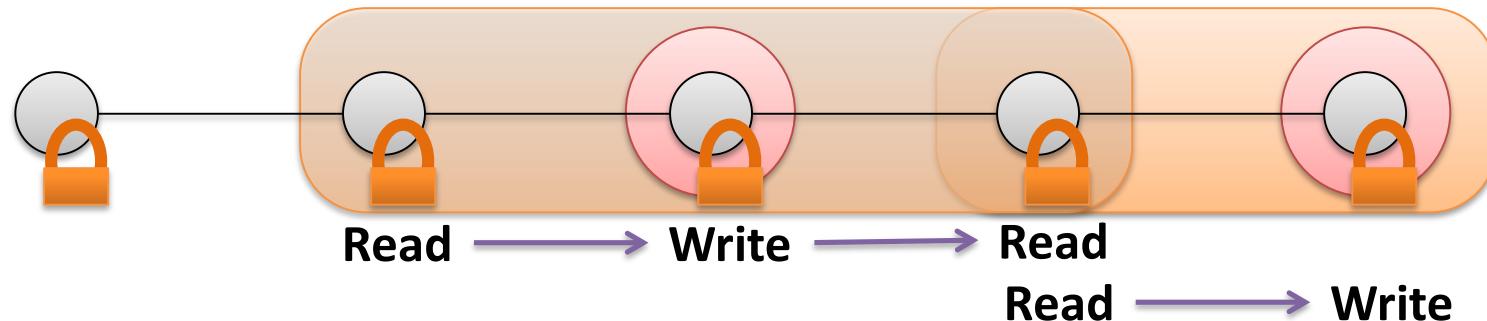
(b) Consistency Models

Consistency Through R/W Locks

- Read/Write locks:
 - Full Consistency

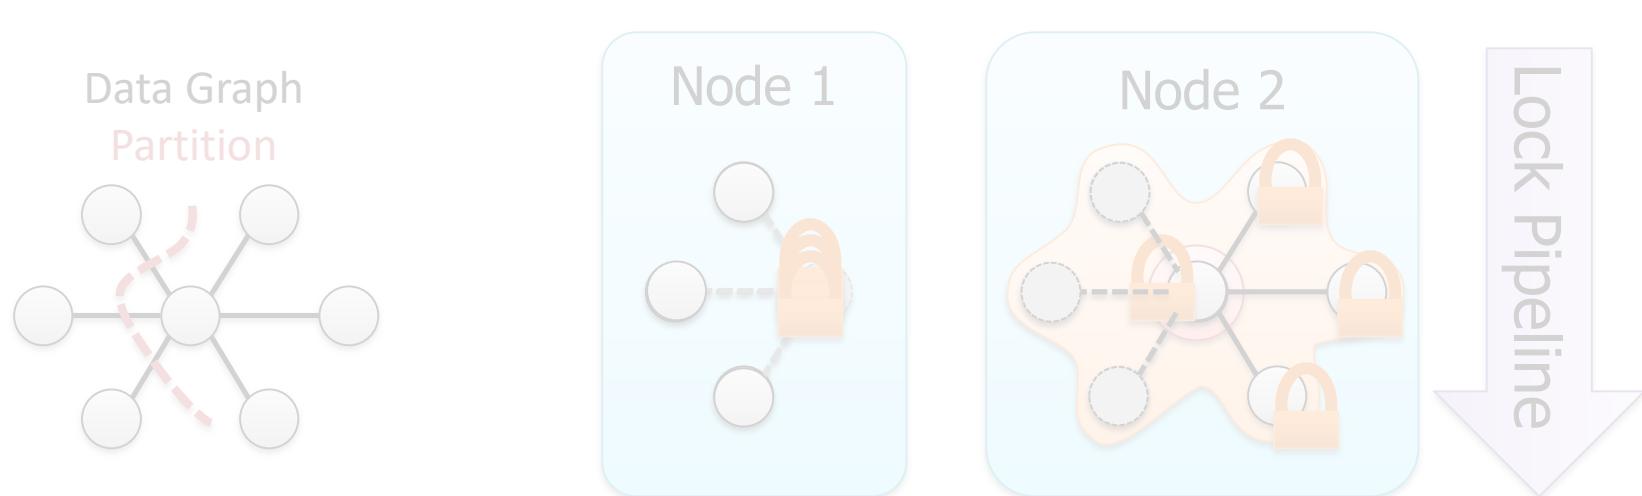


- Edge Consistency



Consistency Through R/W Locks

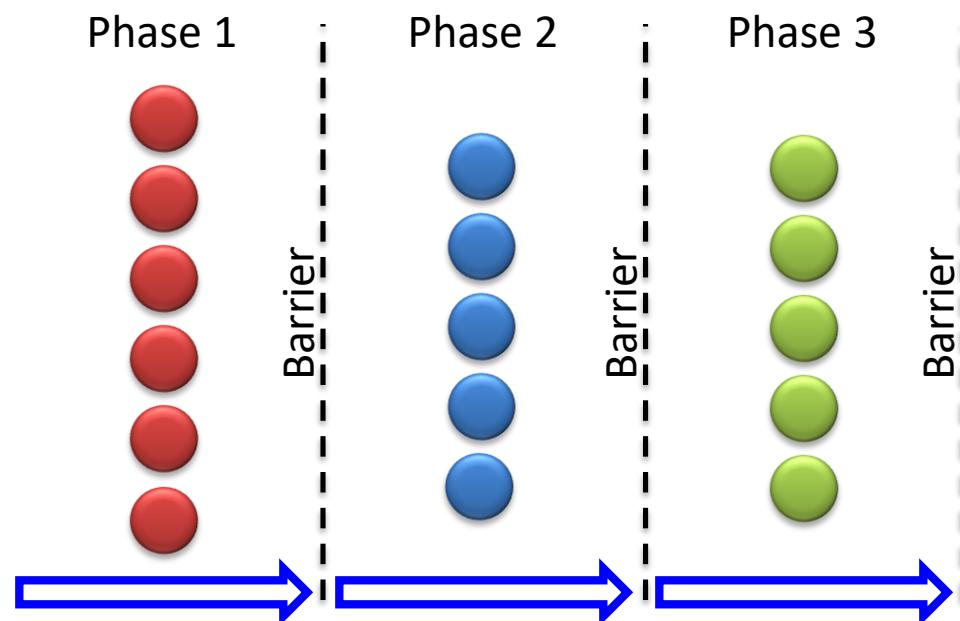
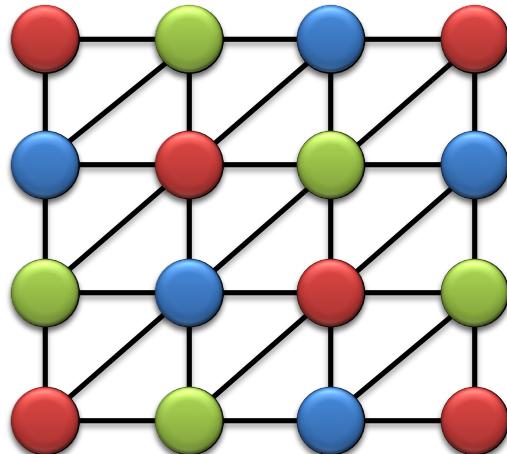
- Multicore Setting: Pthread R/W Locks
- Distributed Setting: *Distributed Locking*
 - Prefetch Locks and Data



- Allow computation to proceed while locks/data are requested.

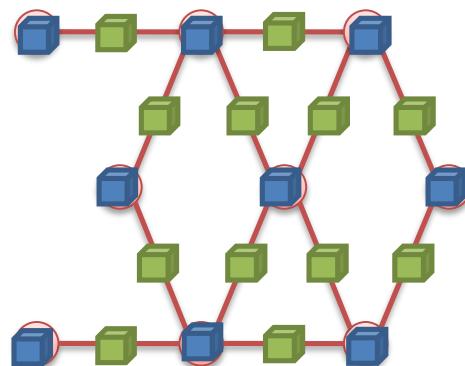
Consistency Through Scheduling

- Edge Consistency Model:
 - Two vertices can be **Updated simultaneously** if they do not share an edge.
- Graph Coloring:
 - Two vertices can be assigned the same color if they do not share an edge.

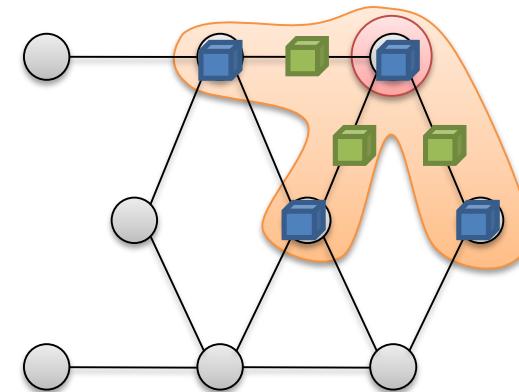


The GraphLab Framework

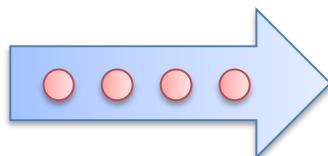
Graph Based
Data Representation



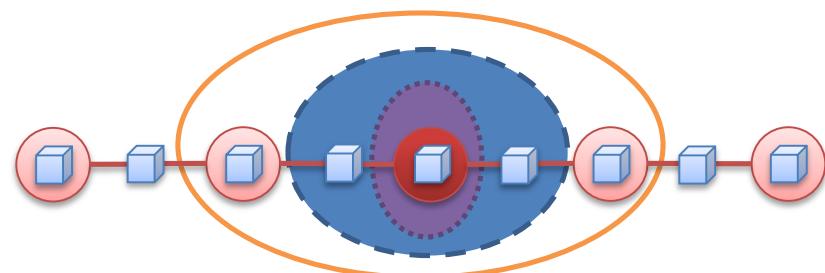
Update Functions
User Computation



Scheduler



Consistency Model



Algorithms Implemented

- PageRank
- Loopy Belief Propagation
- Gibbs Sampling
- CoEM
- Graphical Model Parameter Learning
- Probabilistic Matrix/Tensor Factorization
- Alternating Least Squares
- Lasso with Sparse Features
- Support Vector Machines with Sparse Features
- Label-Propagation
- ...

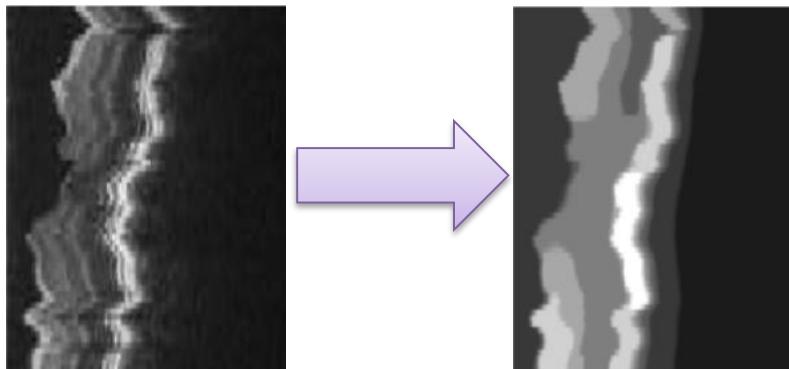
Shared Memory Experiments

Shared Memory Setting

16 Core Workstation

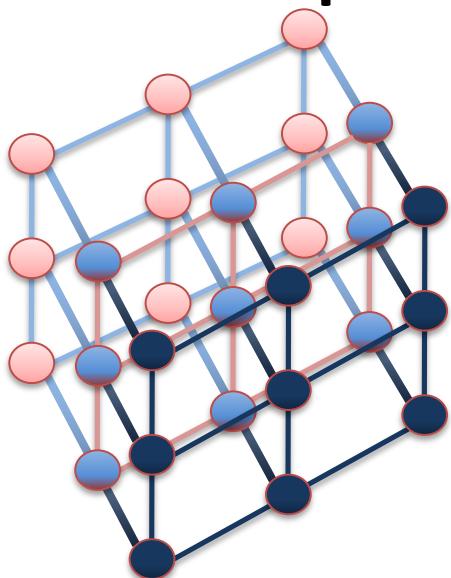
Loopy Belief Propagation

3D retinal image denoising



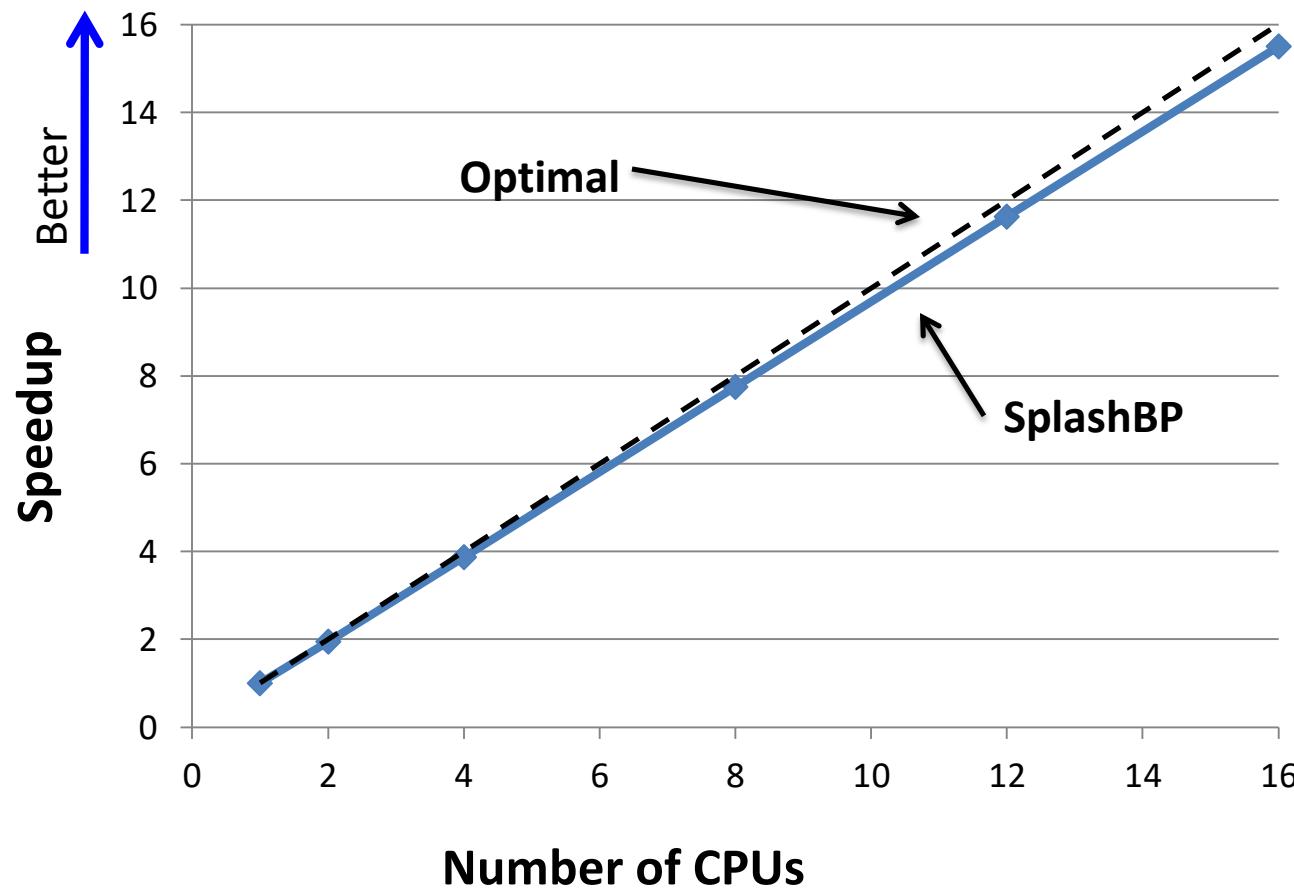
Vertices: 1 Million
Edges: 3 Million

Data Graph



Update Function:
Loopy BP Update Equation
Scheduler:
Approximate Priority
Consistency Model:
Edge Consistency

Loopy Belief Propagation



15.5x speedup

CoEM (Rosie Jones, 2005)

Named Entity Recognition Task

Is “Dog” an animal?

Is “Catalina” a place?

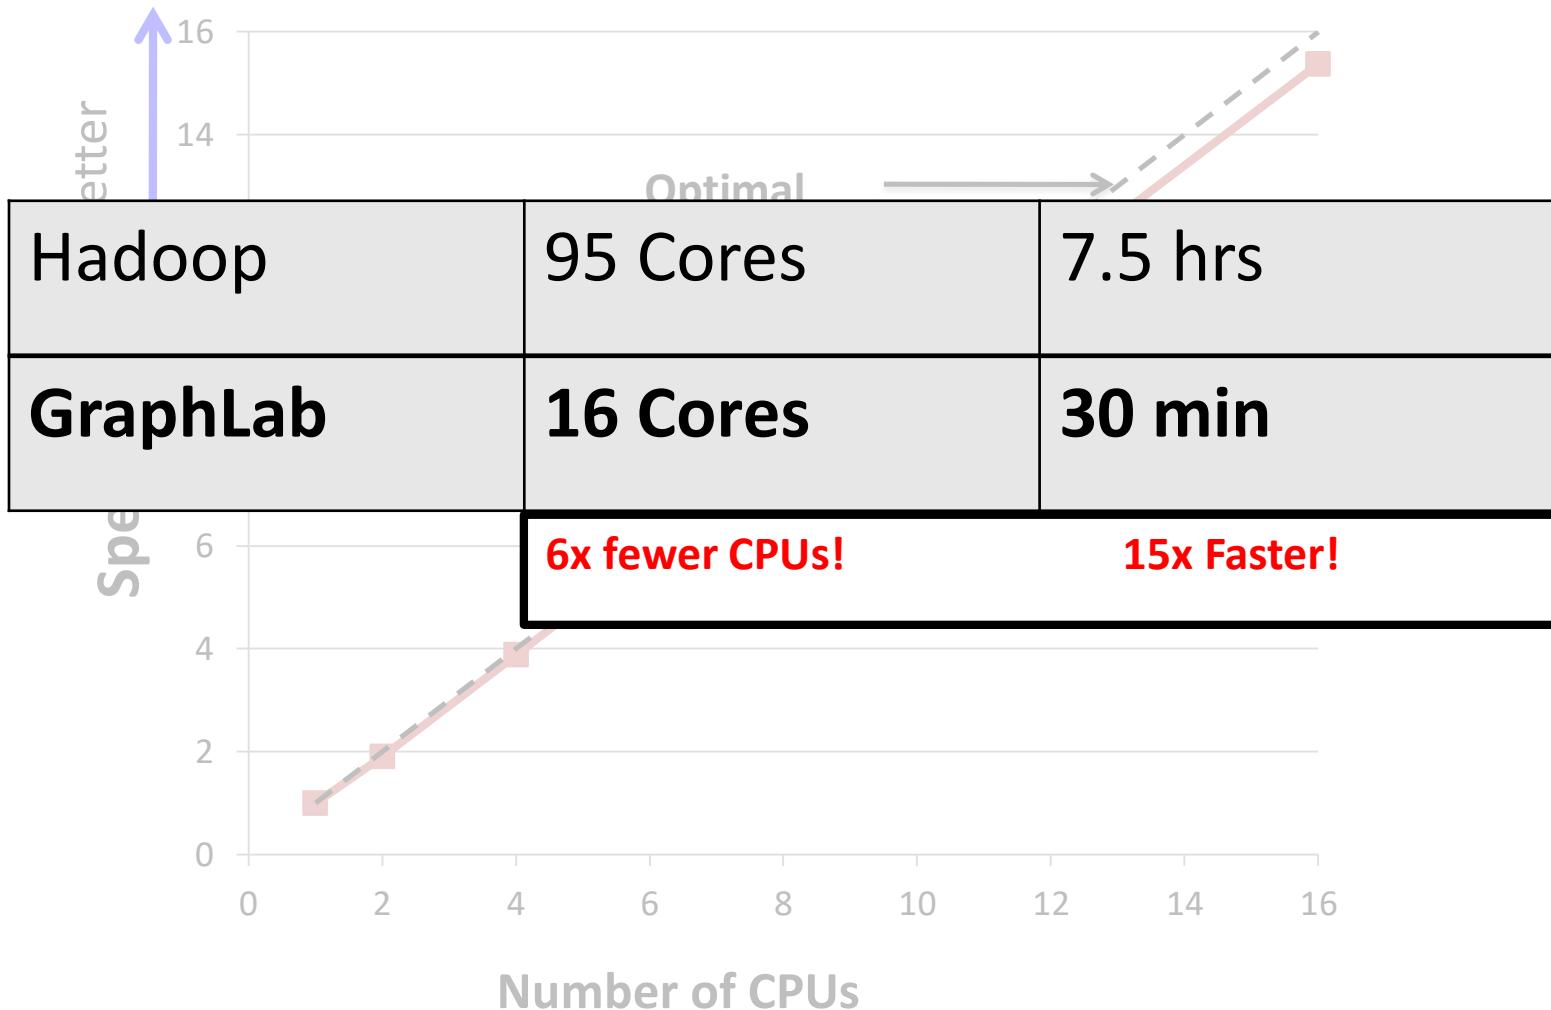
Vertices: 2 Million

Edges: 200 Million



Hadoop	95 Cores	7.5 hrs
--------	----------	---------

CoEM (Rosie Jones, 2005)



Experiments

Amazon EC2
High-Performance Nodes

Video Cosegmentation



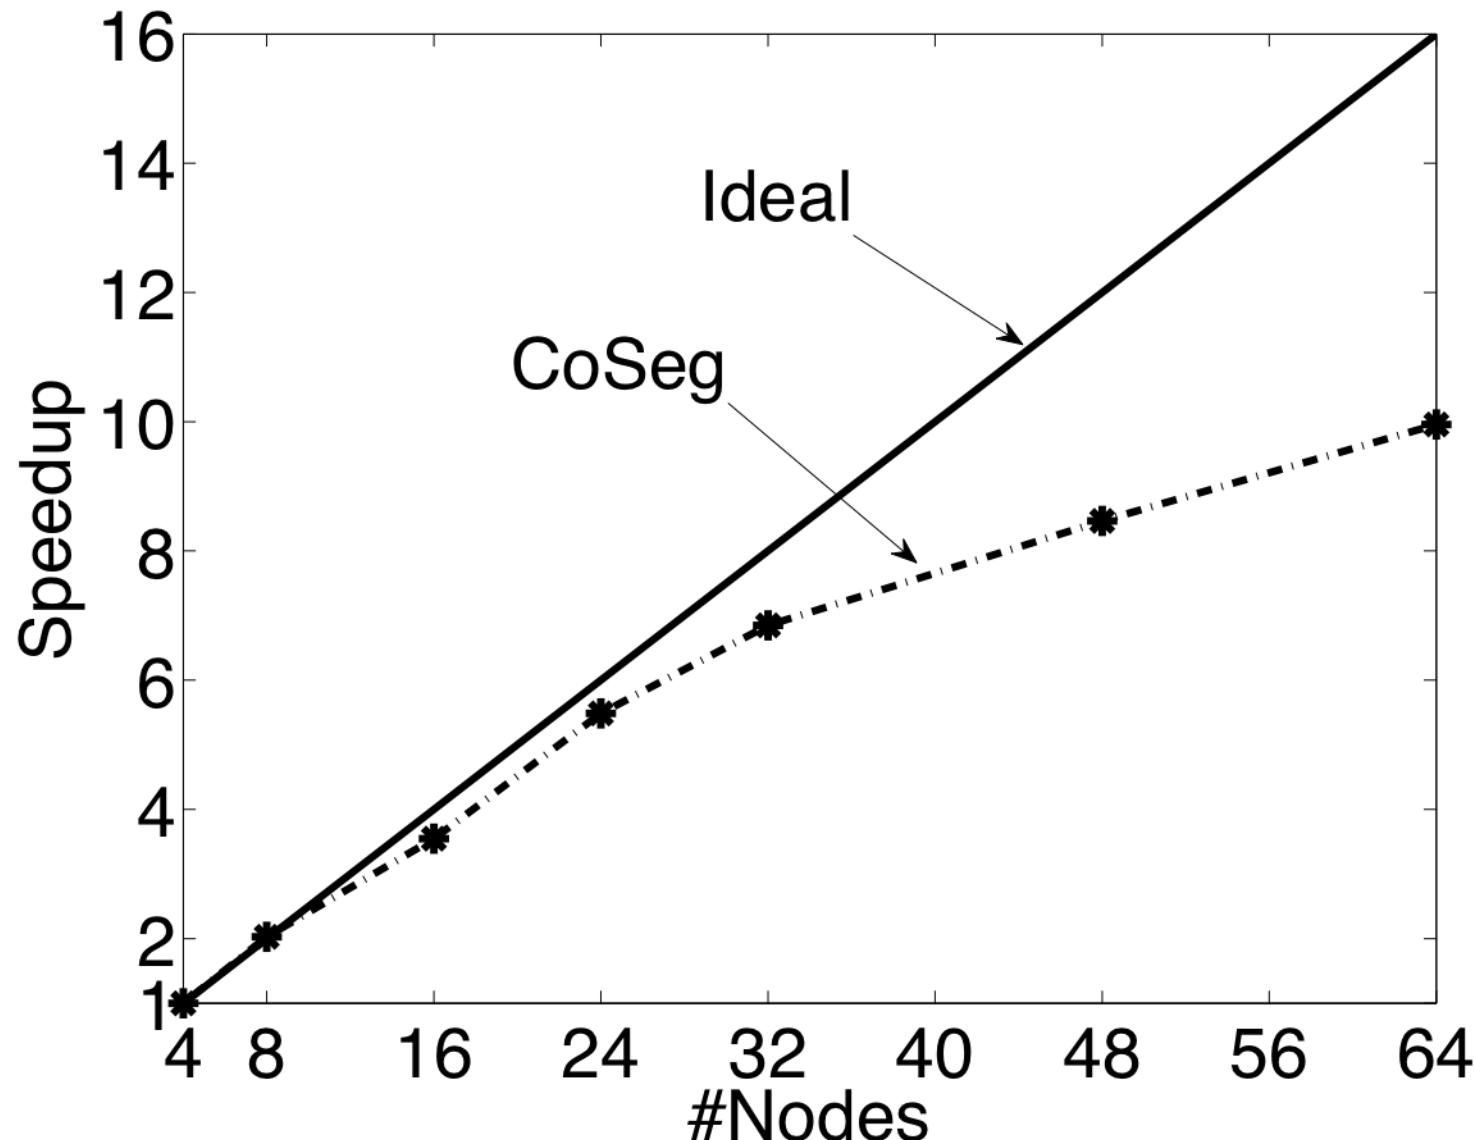
Segments mean the same



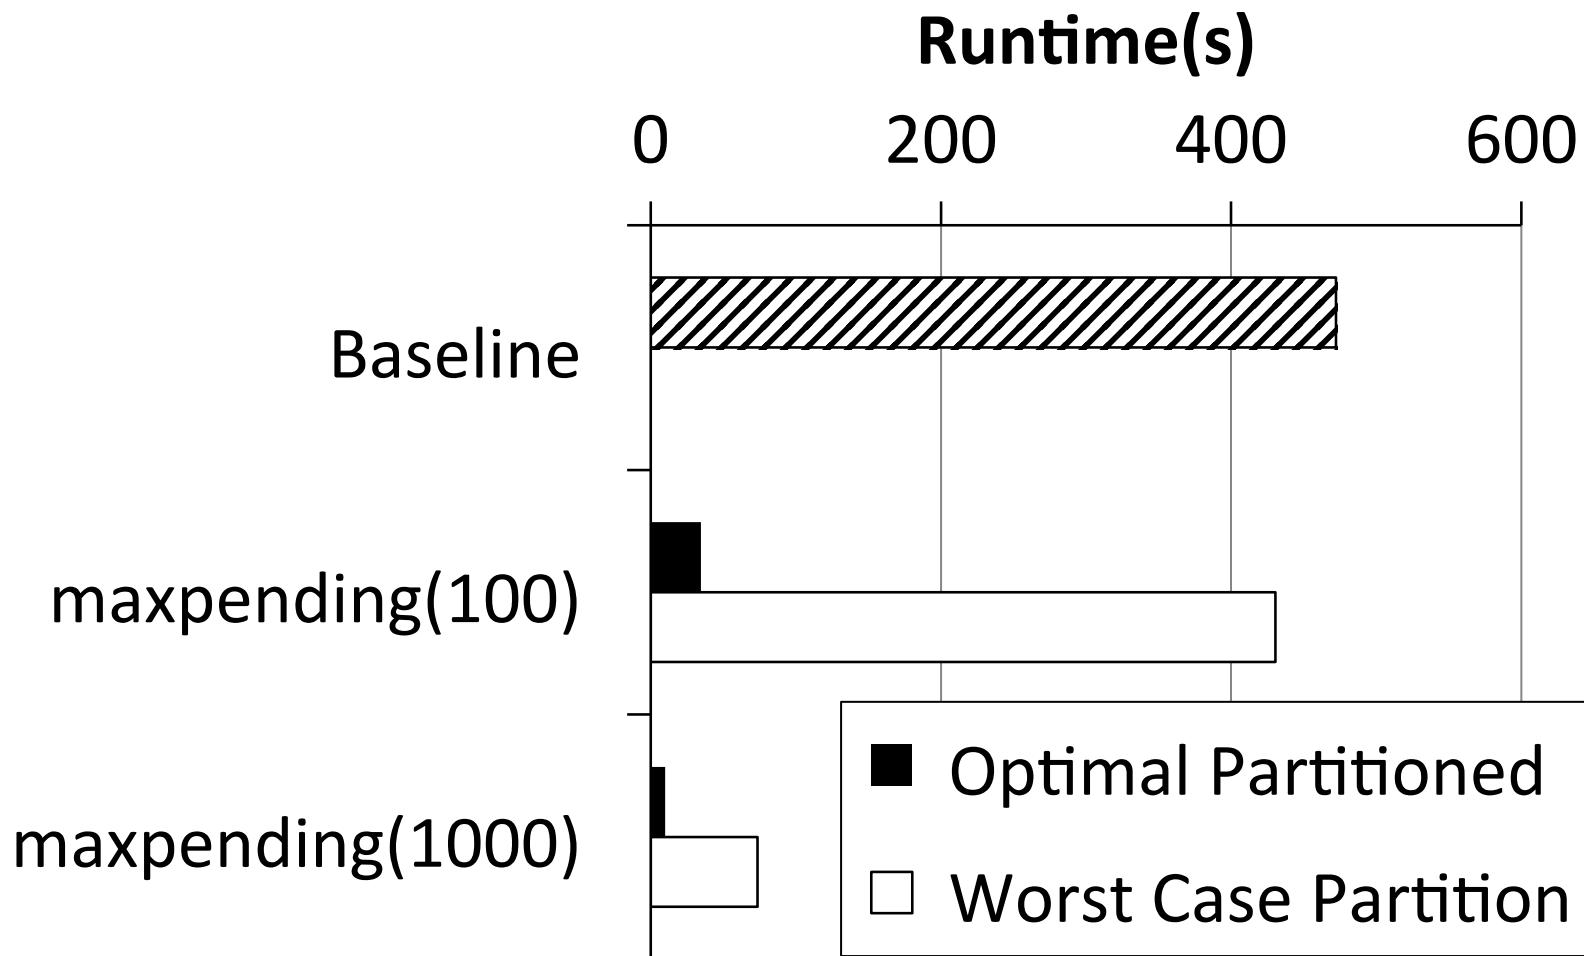
Gaussian EM clustering + BP on 3D grid

Model: 10.5 million nodes, 31 million edges

Video Coseg. Speedups

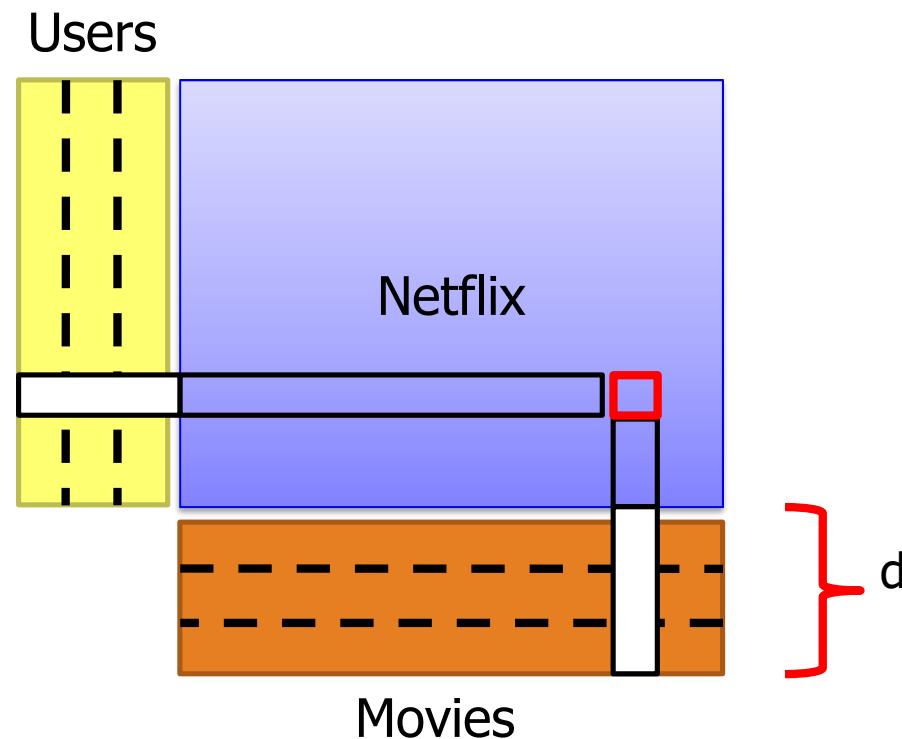


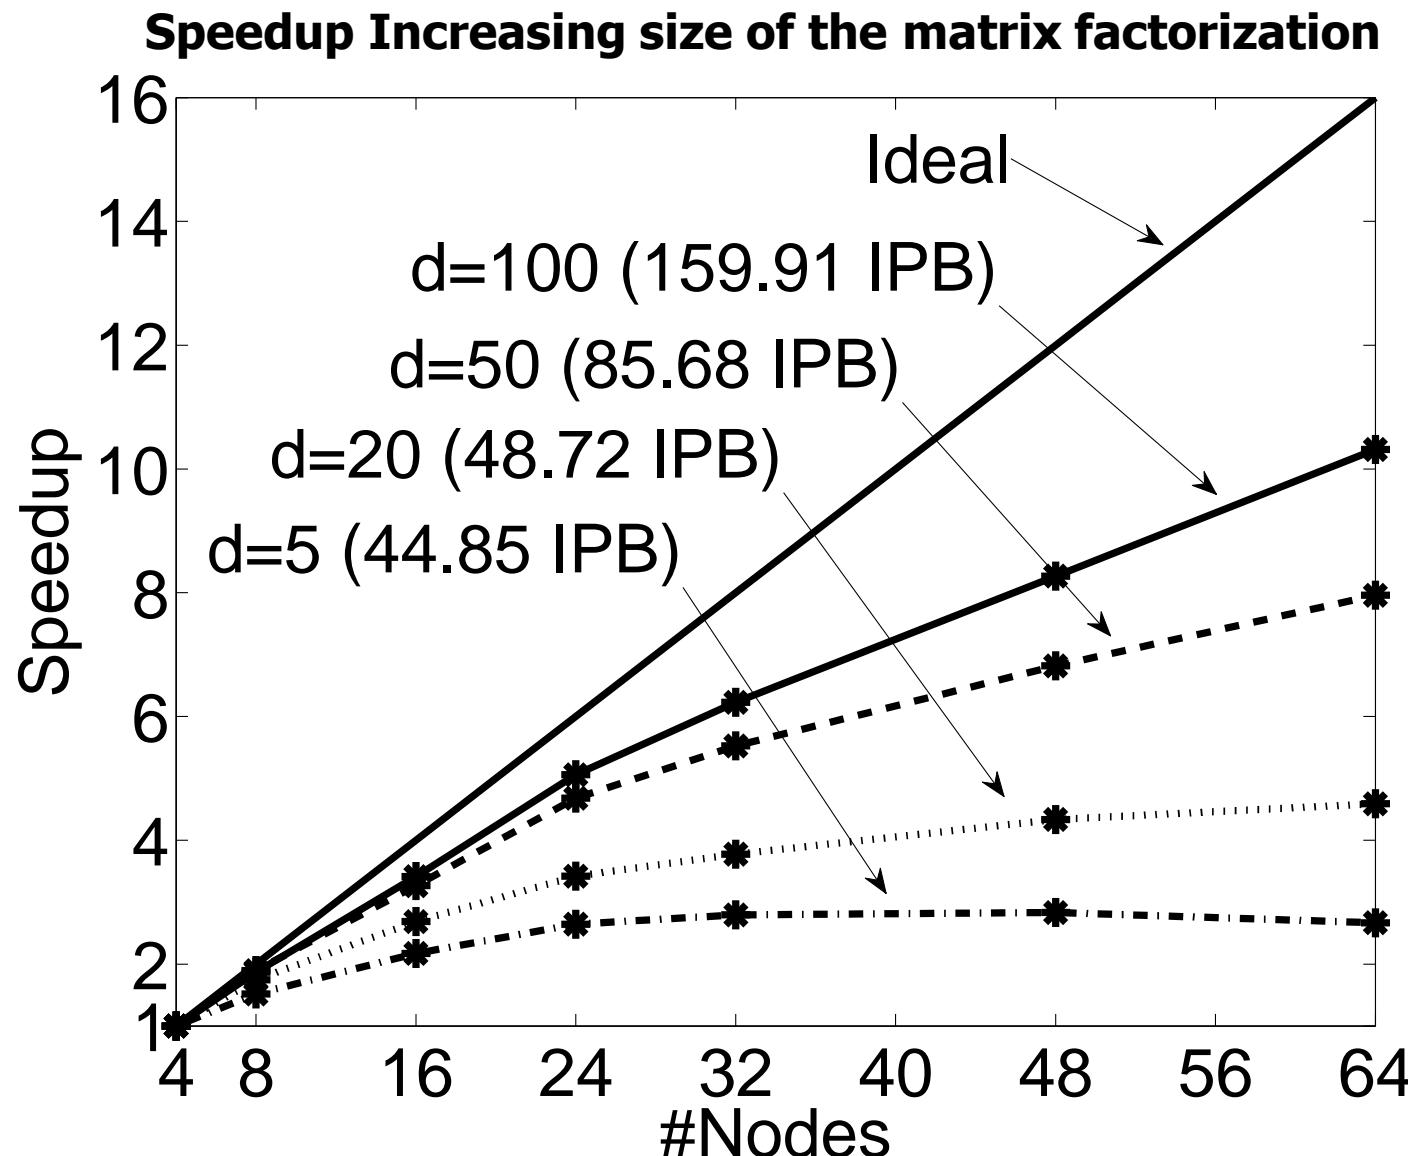
Prefetching Data & Locks



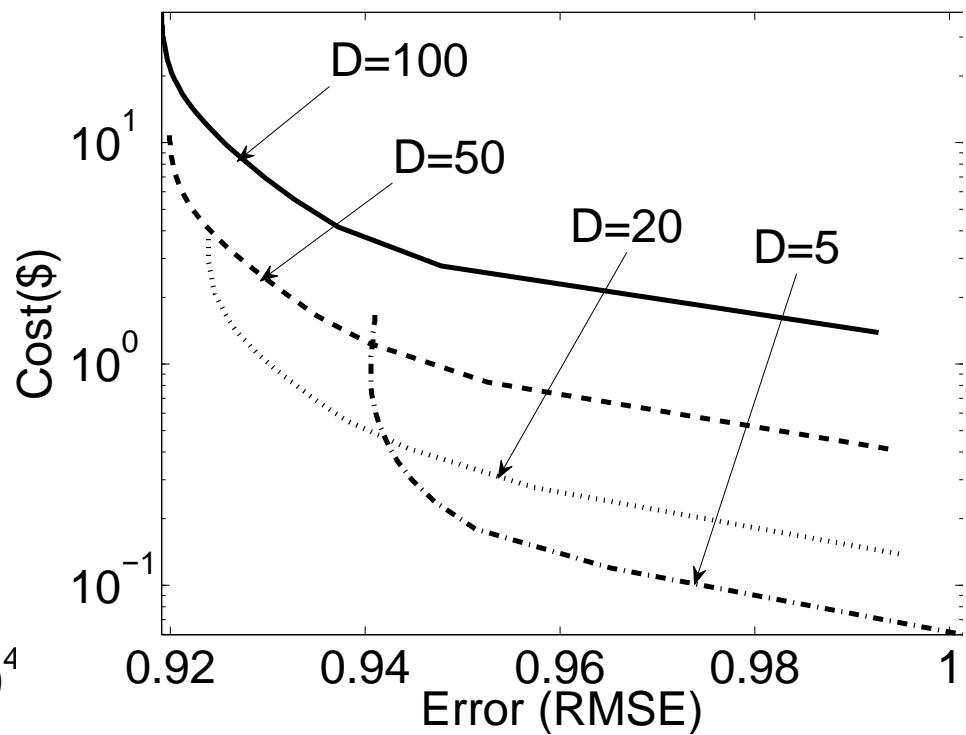
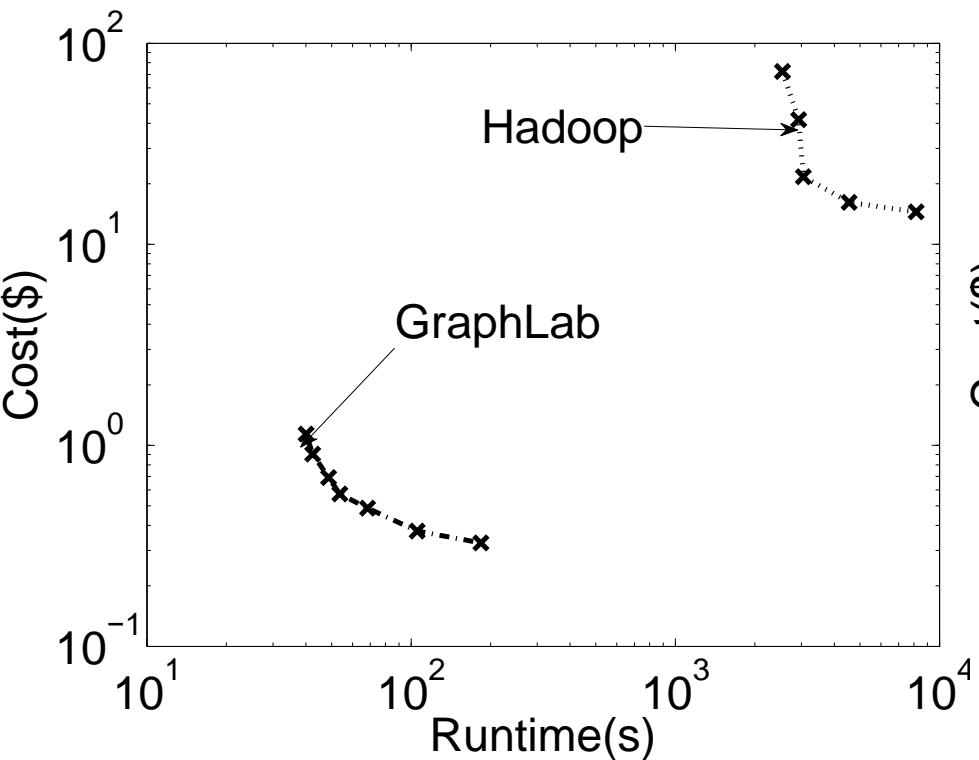
Matrix Factorization

- **Netflix Collaborative Filtering**
 - Alternating Least Squares Matrix Factorization
- Model: 0.5 million nodes, 99 million edges**





The Cost of Hadoop



Summary

- An abstraction tailored to Machine Learning
 - Targets Graph-Parallel Algorithms
- Naturally expresses
 - Data/computational dependencies
 - Dynamic iterative computation
- Simplifies parallel algorithm design
- Automatically ensures data consistency
- Achieves state-of-the-art parallel performance on a variety of problems

Checkout GraphLab

Documentation... Code... Tutorials...

<http://graphlab.org>

Questions & Feedback

jegonzal@cs.cmu.edu

Current/Future Work

- Out-of-core Storage
- Hadoop/HDFS Integration
 - Graph Construction
 - Graph Storage
 - Launching GraphLab from Hadoop
 - Fault Tolerance through HDFS Checkpoints
- Sub-scope parallelism
 - Address the challenge of very high degree nodes
- Improved graph partitioning
- Support for dynamic graph structure